


Homework 7 (due Friday Nov 11)

introduce Q1 today (will introduce Q2 on Thu)

This assignment asks you to create multi-layer networks that learn to classify images of clothing. Fashion-MNIST is a dataset of clothing images consisting of a training set of 60,000 examples and a test set of 10,000 examples. Each example is a 28x28 grayscale image, associated with a label from 10 classes.

see `Homework7_2022.ipynb`

pullover




shirt




sneaker



shirt



dress



coat




trousers



t-shirt



trousers



sneaker



dress




sneaker



boot



dress



sneaker




pullover




sandal



bag



pullover




sandal



coat




pullover



sneaker



dress




bag




bag



trousers



coat



bag



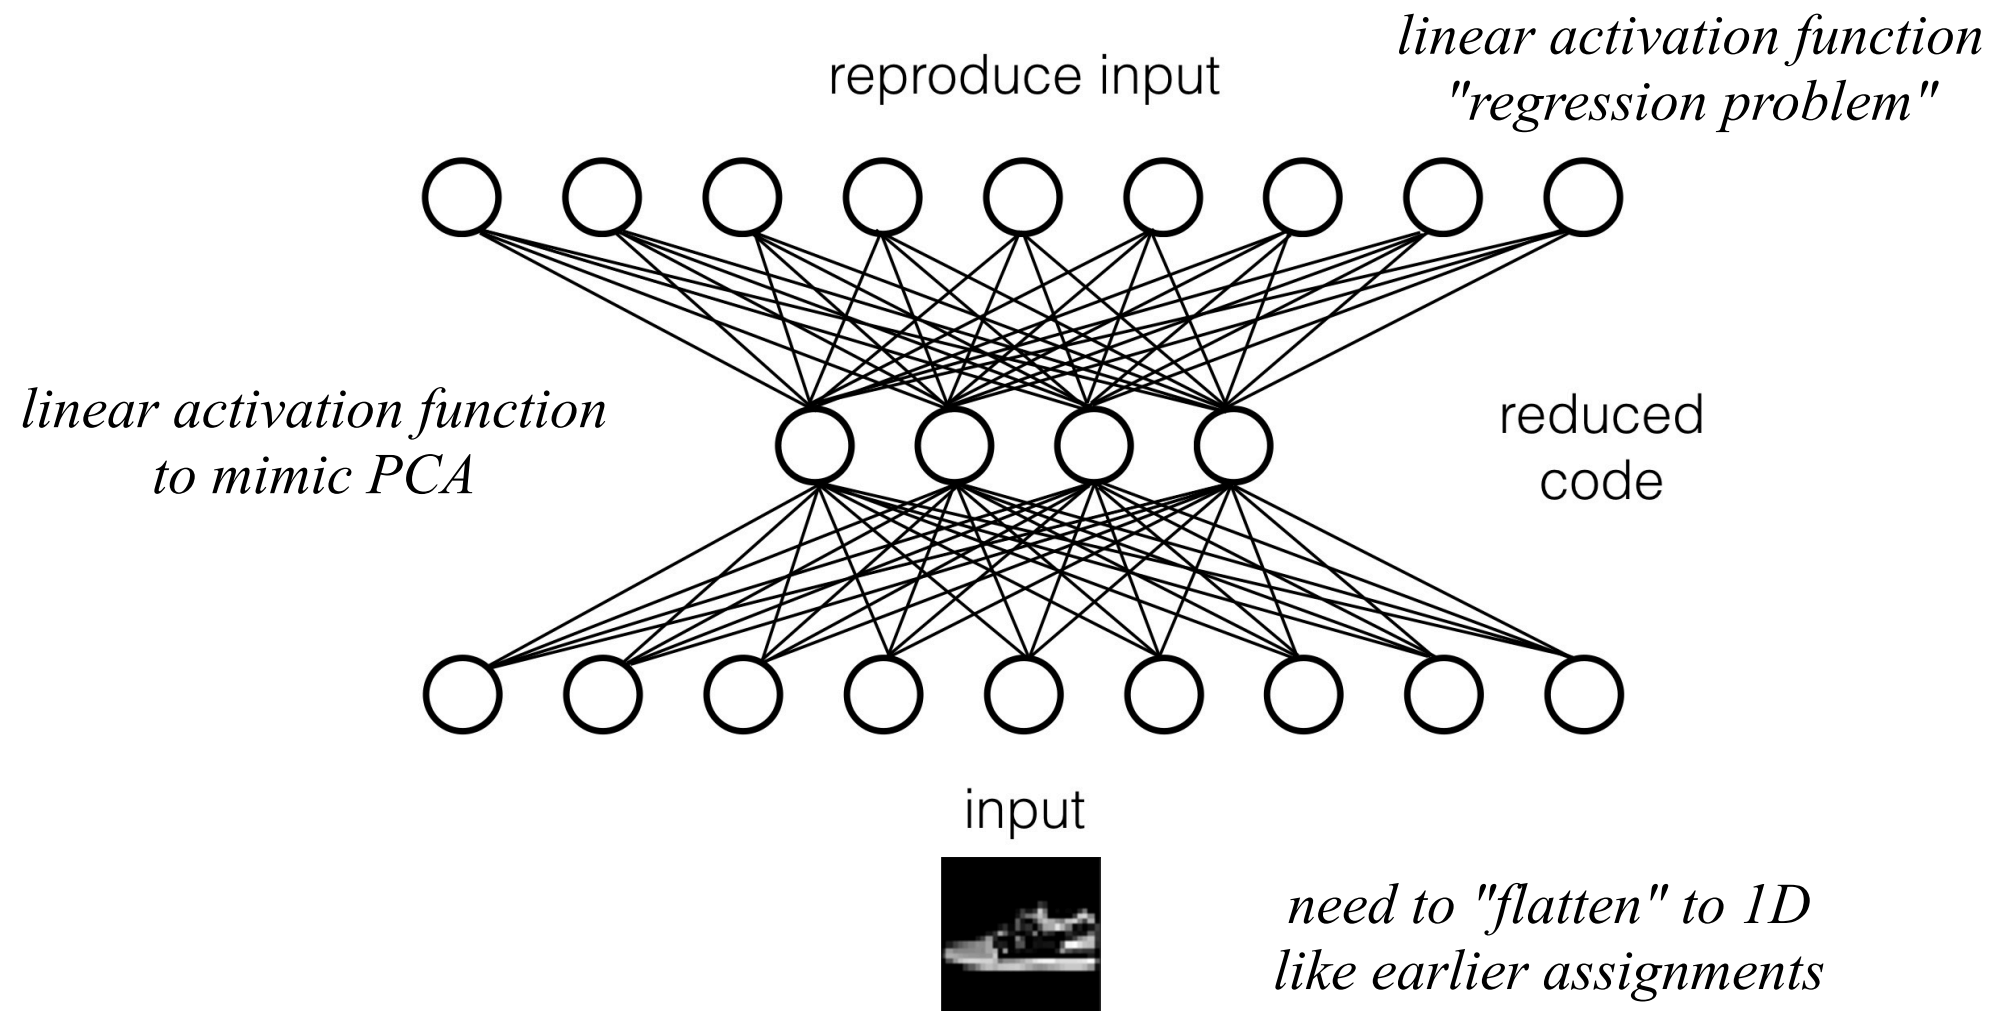
dress



Homework 7 (due Fri Nov 11)

Q1a (5 points). Create an autoassociator network that produces a lower-dimensional representation of the Fashion-MNIST images on its hidden layer. It will use a “flattened” (one-dimensional) version of these images ...


train with 1/2 of the Fashion-MNIST images



Homework 7 (due Thu Nov 11)

... After you train your autoassociator, you will need to create a reduced-dimensionality versions of the remaining 1/2 of the Fashion-MNIST training set. ...

remaining 1/2 of the Fashion-MNIST images




Homework 7 (due Thu Nov 11)

... After you train your autoassociator, you will need to create a reduced-dimensionality versions of the remaining 1/2 of the Fashion-MNIST training set. ...

... The easiest way to do this is to use the weights and biases from the trained autoassociator to calculate the activation values on the hidden layer

```
# get weights and biases from the aa network  
aaW = aanetwork.layers[0].get_weights()[0]  
aaB = aanetwork.layers[0].get_weights()[1]
```



Homework 7 (due Thu Nov 11)

Q1b (3 points). Now, use these lower-dimensional (PCA-like) representations of the clothing images (remaining 1/2 of the original set) as the inputs to train a densely-connected multi-layered neural network that learns, via backpropagation, to classify the images as one of the 10 types of clothing.


```
# get weights and biases from the aa network  
aaW = aanetwork.layers[0].get_weights()[0]  
aaB = aanetwork.layers[0].get_weights()[1]
```

using hidden layer activations from the autoassociator
as input to a multi-layer classification network

autoassociator

from Week10a


once the reduced code is learned,
that code can be used as input to a backprop network



autoassociator


from Week10a

(instead of "fixing these weights", just use the weights and biases to calculate the activations, which will then be used as inputs to a trained network)



use autoassociator representations in feedforward net

from Week10a




autoassociators and principal components analysis (PCA)


*type of neural
network*

statistical technique


autoassociator

using supervised learning to do unsupervised learning







784 output units



784 input units



represent images in
784-dimensional space
using fewer
hidden unit dimensions




output images



with 100 hidden units



input images




output images




with 25 hidden units



input images




output images




with 10 hidden units



input images




output images



with 5 hidden units



input images






What kinds of representations does the autoassociator learn?

Does it try to learn the “parts” that make up different letters?

The autoassociator learns “holistic” building blocks.

assuming 10 hidden units



images produced by weights



training
faces



reproduced



input


training
faces



eigenfaces




face space




20-dimensional face
space rather than a
60000-dimensional
pixel image space

autoassociator


once the reduced code is learned,
that code can be used as input to a backprop network



autoassociator



use autoassociator representations in feedforward net



Homework 7 (due Thu Nov 11)

Q2 (10 points). Create a convolutional neural network that classifies the Fashion-MNIST images. ...

(how to do this will be described today and Thu)

Programming Deep Convolutional Neural Networks using Keras / Tensorflow

```
# this initializes a blank Sequential network  
network = models.Sequential()
```

```
# this initializes a blank Sequential network  
network = models.Sequential()
```

```
# first convolutional layer  
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```




`layers.Conv2D` is a type of layer
(before we used `layers.Dense`)

```
# this initializes a blank Sequential network
network = models.Sequential()
```

```
# first convolutional layer
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```

number of
feature maps



```
network = models.Sequential()
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))

for layer in network.layers:
    print('layer name : {} | input shape : {} | output shape : {}'.
          format(layer.name, layer.input.shape, layer.output.shape))
```


```
layer name : conv2d_1 | input shape : (None, 28, 28, 1) | output shape : (None, 26, 26, 16)
```

```
# this initializes a blank Sequential network
network = models.Sequential()
```

```
# first convolutional layer
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```



dimensions of
the convolution



this initializes a blank Sequential network

```
network = models.Sequential()
```

first convolutional layer

```
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```




activation
function


```
# this initializes a blank Sequential network  
network = models.Sequential()
```

```
# first convolutional layer
```

```
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```


first layer requires
input shape



```
network = models.Sequential()
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))

for layer in network.layers:
    print('layer name : {} | input shape : {} | output shape : {}'.
          format(layer.name, layer.input.shape, layer.output.shape))
```


layer name : conv2d_1 | input shape : (None, 28, 28, 1) | output shape : (None, 26, 26, 16)



```
network = models.Sequential()  
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```

```
for layer in network.layers:  
    print('layer name : {} | input shape : {} | output shape : {}'.  
          format(layer.name, layer.input.shape, layer.output.shape))
```


layer name : conv2d_1 | input shape : (None, 28, 28, 1) | output shape : (None, 26, 26, 16)



```
network = models.Sequential()
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))

for layer in network.layers:
    print('layer name : {} | input shape : {} | output shape : {}'.
          format(layer.name, layer.input.shape, layer.output.shape))
```

layer name : conv2d_1 | input shape : (None, 28, 28, 1) | output shape : (None, 26, 26, 16)



```
# this initializes a blank Sequential network  
network = models.Sequential()
```

```
# first convolutional layer  
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```




required with
Conv2D layer

*what is this
third dimension?*

```
# load fashion mnist dataset
```

```
from tensorflow.keras.datasets import fashion_mnist
```

```
(train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
```



*remember that we reshape this from a matrix
to a vector when training Dense networks*

now we're keeping the 2D structure, but ...

```
# load fashion mnist dataset
from tensorflow.keras.datasets import fashion_mnist

(train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()

print(train_images.shape)
(60000, 28, 28)
```

```
# load fashion mnist dataset
from tensorflow.keras.datasets import fashion_mnist

(train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()

print(train_images.shape)
(60000, 28, 28)

network = models.Sequential()
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))

for layer in network.layers:
    print('layer name : {} | input shape : {} | output shape : {}'.
          format(layer.name, layer.input.shape, layer.output.shape))

layer name : conv2d_1 | input shape : (?, 28, 28, 1) | output shape : (?, 26, 26, 16)
```

error without
this
↓


```
# load fashion mnist dataset
from tensorflow.keras.datasets import fashion_mnist

(train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()

print(train_images.shape)
(60000, 28, 28)

network = models.Sequential()
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))

for layer in network.layers:
    print('layer name : {} | input shape : {} | output shape : {}'.
          format(layer.name, layer.input.shape, layer.output.shape))

layer name : conv2d_1 | input shape : (None, 28, 28, 1) | output shape : (None, 26, 26, 16)

need to reshape the train_images (and test_images)
train_images = train_images.reshape((60000, 28, 28, 1))


        (use variables rather than hard-code)
```

```
# this initializes a blank Sequential network
network = models.Sequential()
```

```
# first convolutional layer
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```

↑
, padding='valid', optional parameter
default is 'valid'

with 'valid' it stays
inside previous layer




```
# this initializes a blank Sequential network  
network = models.Sequential()
```

```
# first convolutional layer  
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```

what will be the dimensions of the Conv2D layer?

with 'valid' it stays
inside previous layer



```
network = models.Sequential()
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))

for layer in network.layers:
    print('layer name : {} | input shape : {} | output shape : {}'.
          format(layer.name, layer.input.shape, layer.output.shape))
```

```
layer name : conv2d_1 | input shape : (None, 28, 28, 1) | output shape : (None, 26, 26, 16)
```

```
# this initializes a blank Sequential network
network = models.Sequential()
```


```
# first convolutional layer
```

```
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```

may want to use 'same' in Homework 7

↑
, padding='same', optional parameter default is 'valid'

with 'same' it goes outside previous layer




```
# this initializes a blank Sequential network
network = models.Sequential()
```

may want to use 'same' in Homework 7

```
# first convolutional layer
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```

what will be the dimensions of the Conv2D layer?



```
network = models.Sequential()
network.add(layers.Conv2D(16, (3,3), padding='same', activation='relu',
                          input_shape=(sz,sz,1)))

for layer in network.layers:
    print('layer name : {} | input shape : {} | output shape : {}'.
          format(layer.name, layer.input.shape, layer.output.shape))
```


```
layer name : conv2d_1 | input shape : (None, 28, 28, 1) | output shape : (None, 28, 28, 16)
```

```
# this initializes a blank Sequential network
network = models.Sequential()
```

```
# first convolutional layer
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```

↑
, strides=(1, 1), optional parameter
default is (1,1)

with (1,1) stride,
convolution full covers
previous layer





```
# this initializes a blank Sequential network
network = models.Sequential()
```

```
# first convolutional layer
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```

↑
, strides=(1, 1), optional parameter
default is (1,1)

with (1,1) stride,
convolution full covers
previous layer




```
# this initializes a blank Sequential network
network = models.Sequential()
```

```
# first convolutional layer
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```

↑
, strides=(1, 1), optional parameter
default is (1,1)

with (1,1) stride,
convolution full covers
previous layer




```
# this initializes a blank Sequential network
network = models.Sequential()
```

```
# first convolutional layer
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```

↑
, strides=(1, 1), optional parameter
default is (1,1)

with (1,1) stride,
convolution full covers
previous layer



**usually use
full stride**


```
# this initializes a blank Sequential network  
network = models.Sequential()
```

```
# first convolutional layer
```

```
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```

↑
, strides=(1, 1), optional parameter
default is (1,1)

with (1,1) stride,
convolution full covers
previous layer




```
# this initializes a blank Sequential network
network = models.Sequential()
```

```
# first convolutional layer
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```

↑
, strides=(2, 2), optional parameter
default is (1,1)

with (2,2) stride,
convolution has gaps




```
# this initializes a blank Sequential network
network = models.Sequential()
```

```
# first convolutional layer
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```

↑
, strides=(2, 2), optional parameter
default is (1,1)

with (2,2) stride,
convolution has gaps




```
# this initializes a blank Sequential network
network = models.Sequential()
```

```
# first convolutional layer
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```

↑
, strides=(2, 2), optional parameter
default is (1,1)

with (2,2) stride,
convolution has gaps




```
# this initializes a blank Sequential network
network = models.Sequential()
```

```
# first convolutional layer
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```

↑
, strides=(2, 2), optional parameter
default is (1,1)

with (2,2) stride,
convolution has gaps





```
# this initializes a blank Sequential network
network = models.Sequential()
```

```
# first convolutional layer
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```

↑
, strides=(2, 2), optional parameter
default is (1,1)

with (2,2) stride,
convolution has gaps




```
# this initializes a blank Sequential network
network = models.Sequential()
```

```
# first convolutional layer
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
```

↑
, strides=(2, 2), optional parameter
default is (1,1)

what will be the dimensions of the Conv2D layer?

with (2,2) stride,
convolution has gaps



```
network = models.Sequential()  
network.add(layers.Conv2D(16, (3,3), strides=(2,2), activation='relu',  
                           input_shape=(sz,sz,1)))
```


```
for layer in network.layers:  
    print('layer name : {} | input shape : {} | output shape : {}'.  
          format(layer.name, layer.input.shape, layer.output.shape))
```

```
layer name : conv2d_1 | input shape : (None, 28, 28, 1) | output shape : (None, 13, 13, 16)
```


```
network = models.Sequential()  
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))  
network.add(layers.Conv2D(16, (3,3), activation='relu'))
```

```
for layer in network.layers:  
    print('layer name : {} | input shape : {} | output shape : {}'.  
          format(layer.name, layer.input.shape, layer.output.shape))
```

```
layer name : conv2d_1 | input shape : (None, 28, 28, 1) | output shape : (None, 26, 26, 16)  
layer name : conv2d_2 | input shape : (None, 26, 26, 16) | output shape : (None, 24, 24, 16)
```



```
network = models.Sequential()  
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))  
network.add(layers.Conv2D(16, (3,3), activation='relu'))  
network.add(layers.MaxPooling2D((2,2)))
```




what will be the dimensions
of the max_pooling2d layer?

```
network = models.Sequential()  
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))  
network.add(layers.Conv2D(16, (3,3), activation='relu'))  
network.add(layers.MaxPooling2D((2,2)))
```

```
for layer in network.layers:  
    print('layer name : {} | input shape : {} | output shape : {}'.  
          format(layer.name, layer.input.shape, layer.output.shape))
```

```
layer name : conv2d_1 | input shape : (None, 28, 28, 1) | output shape : (None, 26, 26, 16)  
layer name : conv2d_2 | input shape : (None, 26, 26, 16) | output shape : (None, 24, 24, 16)  
layer name : max_pooling2d_3 | input shape : (None, 24, 24, 16) | output shape : (None, 12, 12, 16)
```



⋮

Max Pooling

Convolutional

Max Pooling

Convolutional

Max Pooling

Convolutional

Image

⋮

Max Pooling

Convolutional

Convolutional

Max Pooling

Convolutional

Convolutional

Image

⋮

Max Pooling

Convolutional

Max Pooling


Convolutional

Convolutional

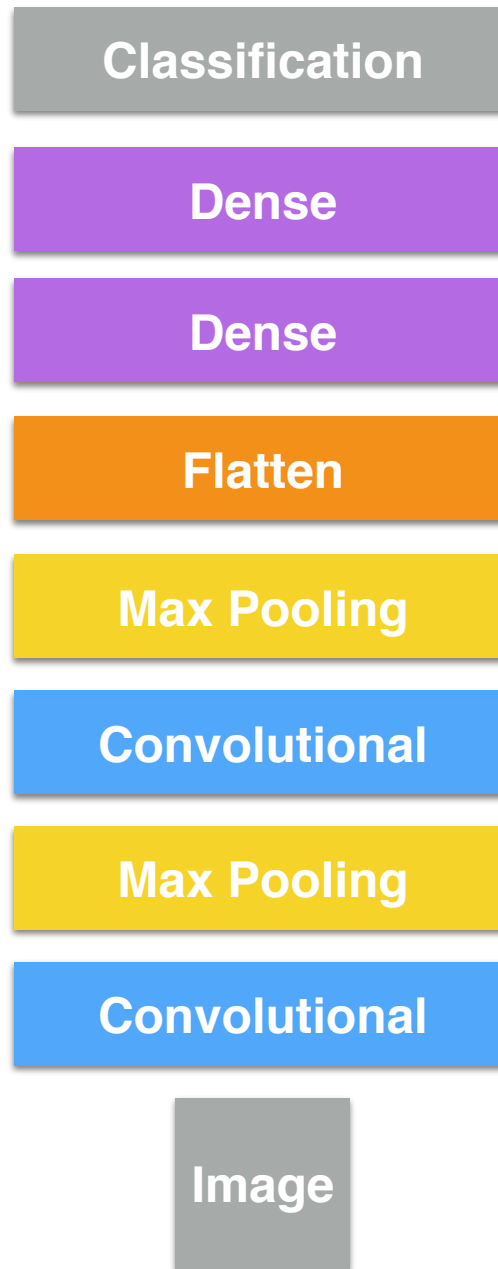
Convolutional

Image

AlexNet



just an example



```
network = models.Sequential()
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
network.add(layers.MaxPooling2D((2,2)))
network.add(layers.Conv2D(16, (3,3), activation='relu'))
network.add(layers.MaxPooling2D((2,2)))
network.add(layers.Flatten())
network.add(layers.Dense(20, activation='relu'))
network.add(layers.Dense(nout, activation='softmax'))

for layer in network.layers:
    print('layer name : {} | input shape : {} | output shape : {}'.
          format(layer.name, layer.input.shape, layer.output.shape))
```

```
layer name : conv2d_1 | input shape : (None, 28, 28, 1) | output shape : (None, 26, 26, 16)
layer name : max_pooling2d_2 | input shape : (None, 26, 26, 16) | output shape : (None, 13, 13, 16)
layer name : conv2d_3 | input shape : (None, 13, 13, 16) | output shape : (None, 11, 11, 16)
layer name : max_pooling2d_4 | input shape : (None, 11, 11, 16) | output shape : (None, 5, 5, 16)
layer name : flatten_5 | input shape : (None, 5, 5, 16) | output shape : (None, 400)
layer name : dense_6 | input shape : (None, 400) | output shape : (None, 20)
layer name : dense_7 | input shape : (None, 20) | output shape : (None, 10)
```

```
print(network.summary())
```

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 26, 26, 16)	160
max_pooling2d_2 (MaxPooling)	(None, 13, 13, 16)	0
conv2d_3 (Conv2D)	(None, 11, 11, 16)	2320
max_pooling2d_4 (MaxPooling)	(None, 5, 5, 16)	0
flatten_5 (Flatten)	(None, 400)	0
dense_6 (Dense)	(None, 20)	8020
dense_7 (Dense)	(None, 10)	210
=====		
Total params: 10,710		
Trainable params: 10,710		
Non-trainable params: 0		

```
network = models.Sequential()
network.add(layers.Conv2D(16, (3,3), activation='relu', input_shape=(sz,sz,1)))
network.add(layers.MaxPooling2D((2,2)))
network.add(layers.Conv2D(16, (3,3), activation='relu'))
network.add(layers.MaxPooling2D((2,2)))
network.add(layers.Flatten())
network.add(layers.Dense(20, activation='relu'))
network.add(layers.Dense(nout, activation='softmax'))

network.compile(optimizer='adam', loss='categorical_crossentropy',
                metrics=['accuracy'])

history = network.fit(train_images, train_labels_onehot, verbose=True,
                    validation_split=.1, epochs=ep, batch_size=128)

test_loss, test_accuracy = network.evaluate(test_images, test_labels_onehot)
```