

Homework 8

Part One now

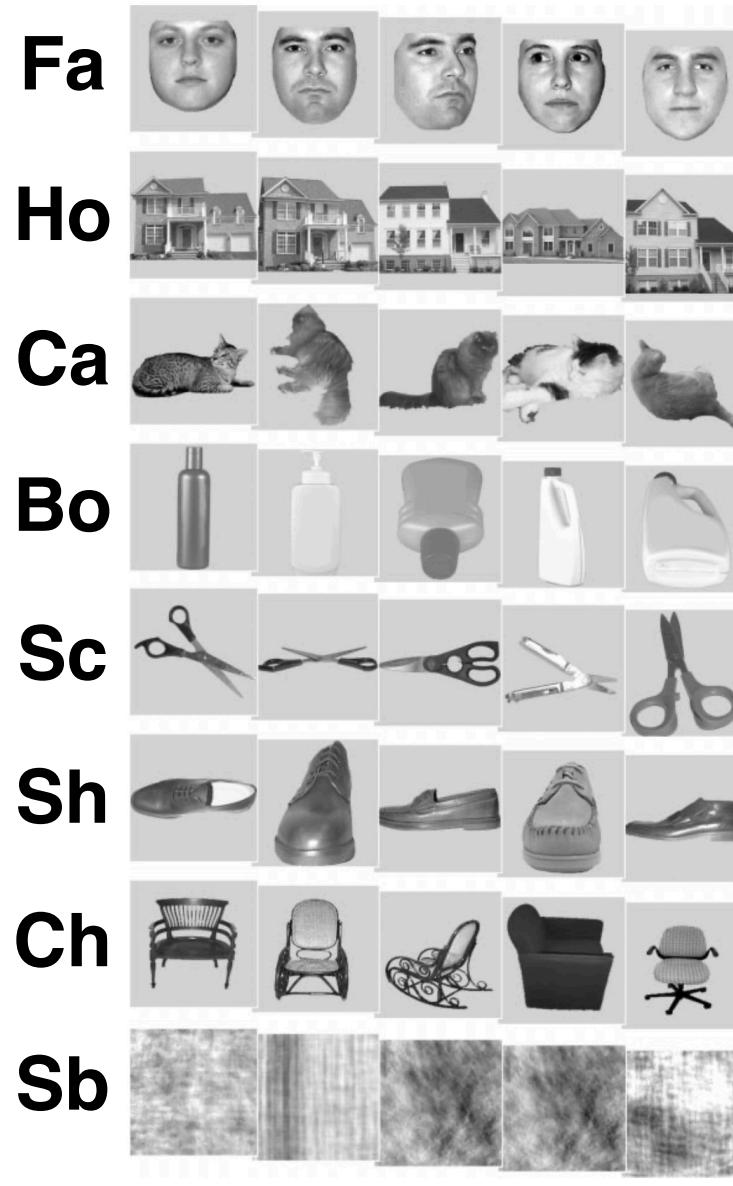
Part Two coming soon

due last day of class (Thu Dec 8)

Part One : reanalysis of Haxby et al. 2001 data

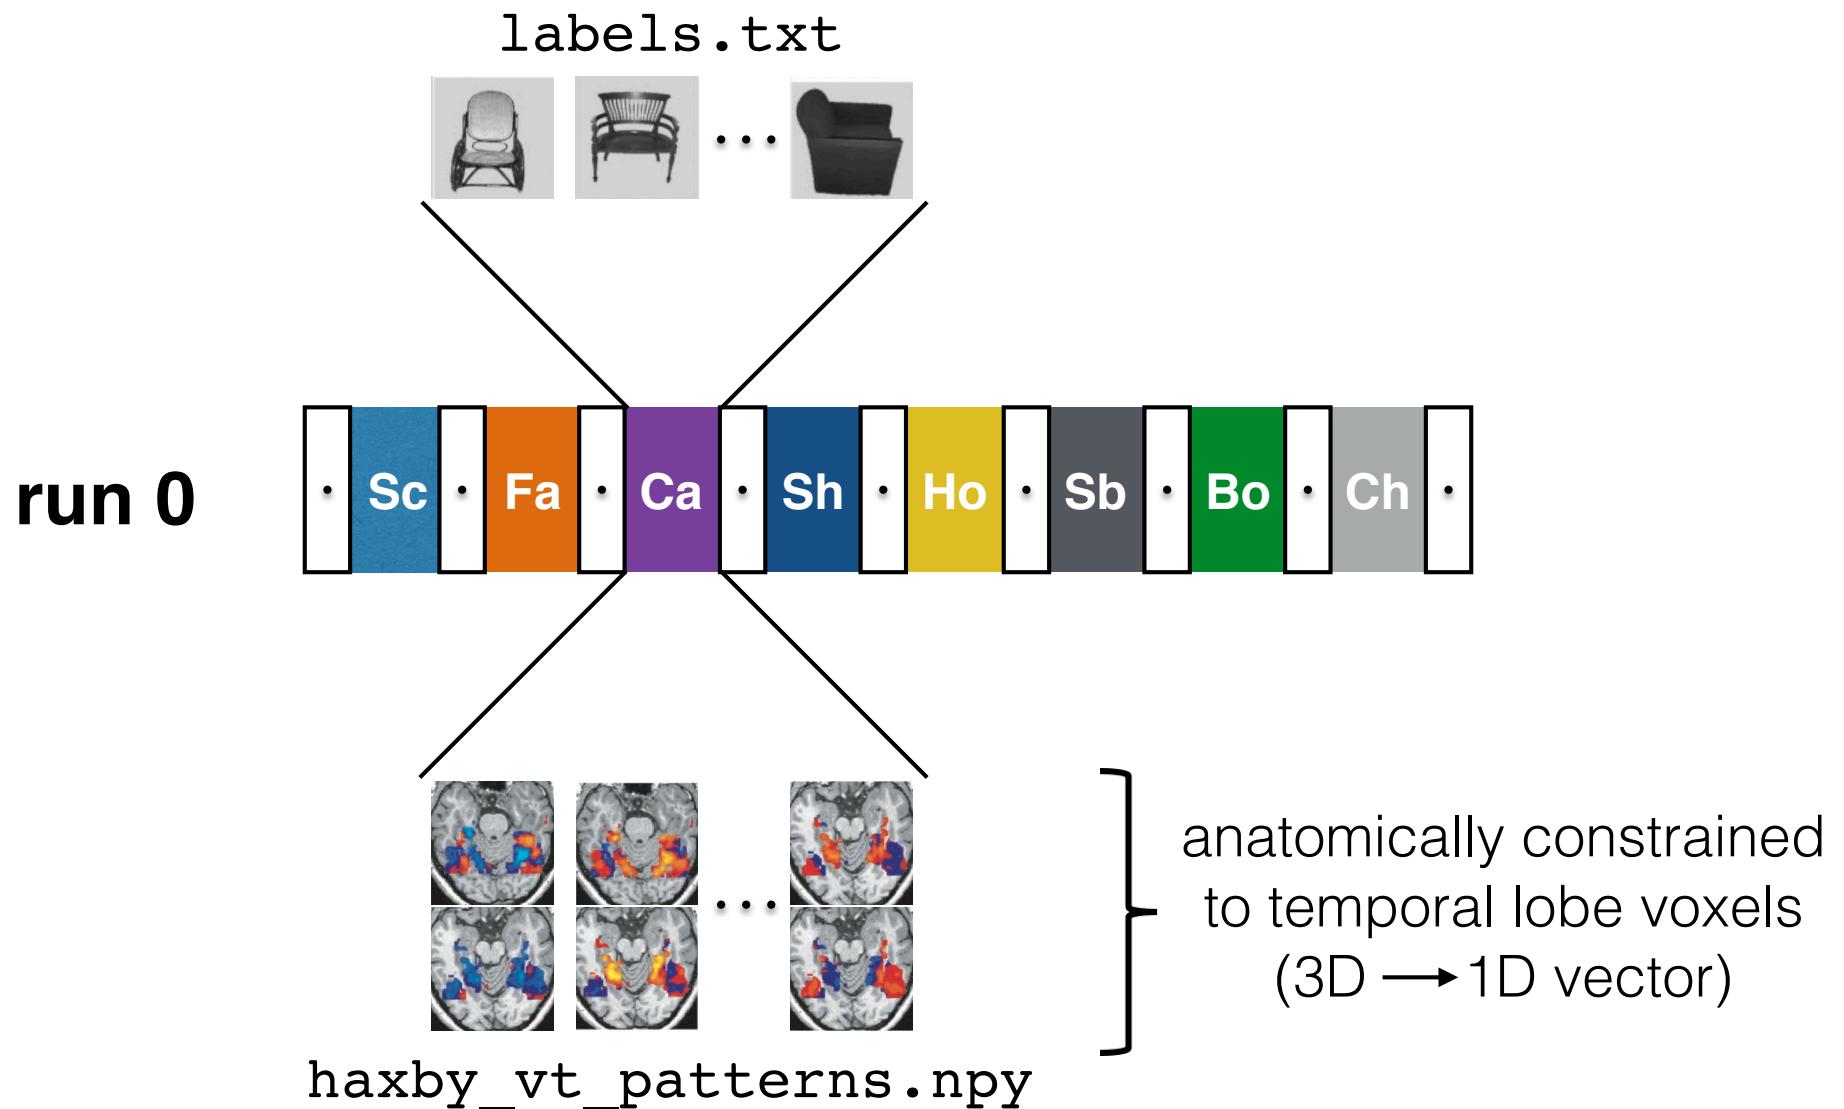
Their original analysis used a simple correlation-based classifier. We will use a simple neural network classifier, trained with backprop.

Haxby et al. 2001 details



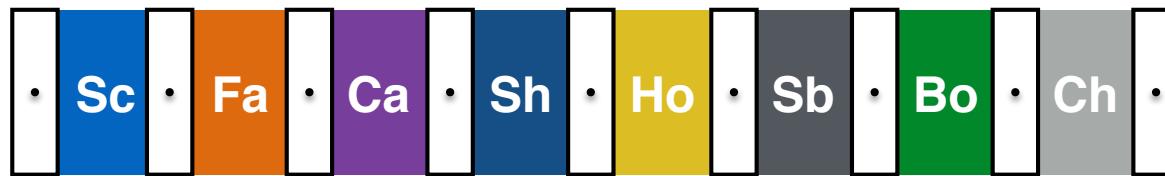
Haxby, J.V., Gobbini, M.I., Furey, M.L., Ishai, A., Schouten, J.L., & Pietrini, P. (2001). Distributed and overlapping representations of faces and objects in ventral temporal cortex. *Science*, 293, 2425-2430.

Haxby et al. 2001 details

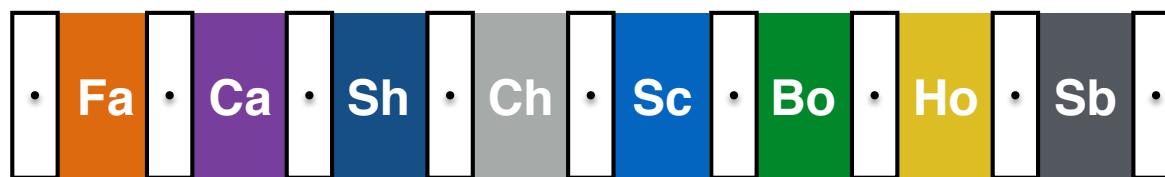


Haxby et al. 2001 details

run 0

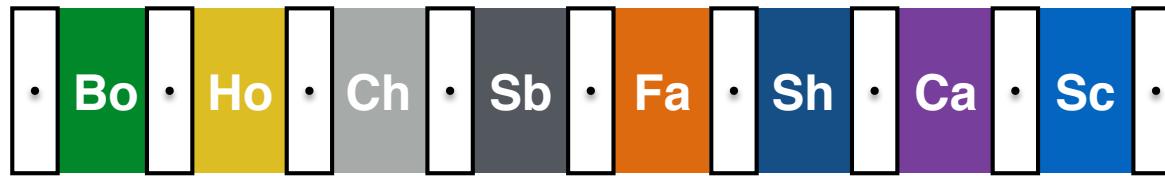


run 1



⋮

run 11



Haxby et al. 2001 details

labels.txt (on Brightspace)

```
labels run
rest 0
scissors 0
rest 0
face 0
.
.
.
scissors 11
scissors 11
scissors 11
scissors 11
rest 11
rest 11
rest 11
rest 11
rest 11
rest 11
```

Homework8.ipynb

```
patterns, targets, runs = prep_haxby_data()
```

```
prep_haxby_data()
```

- opens haxby_vt_patterns.npy
- read in labels.txt

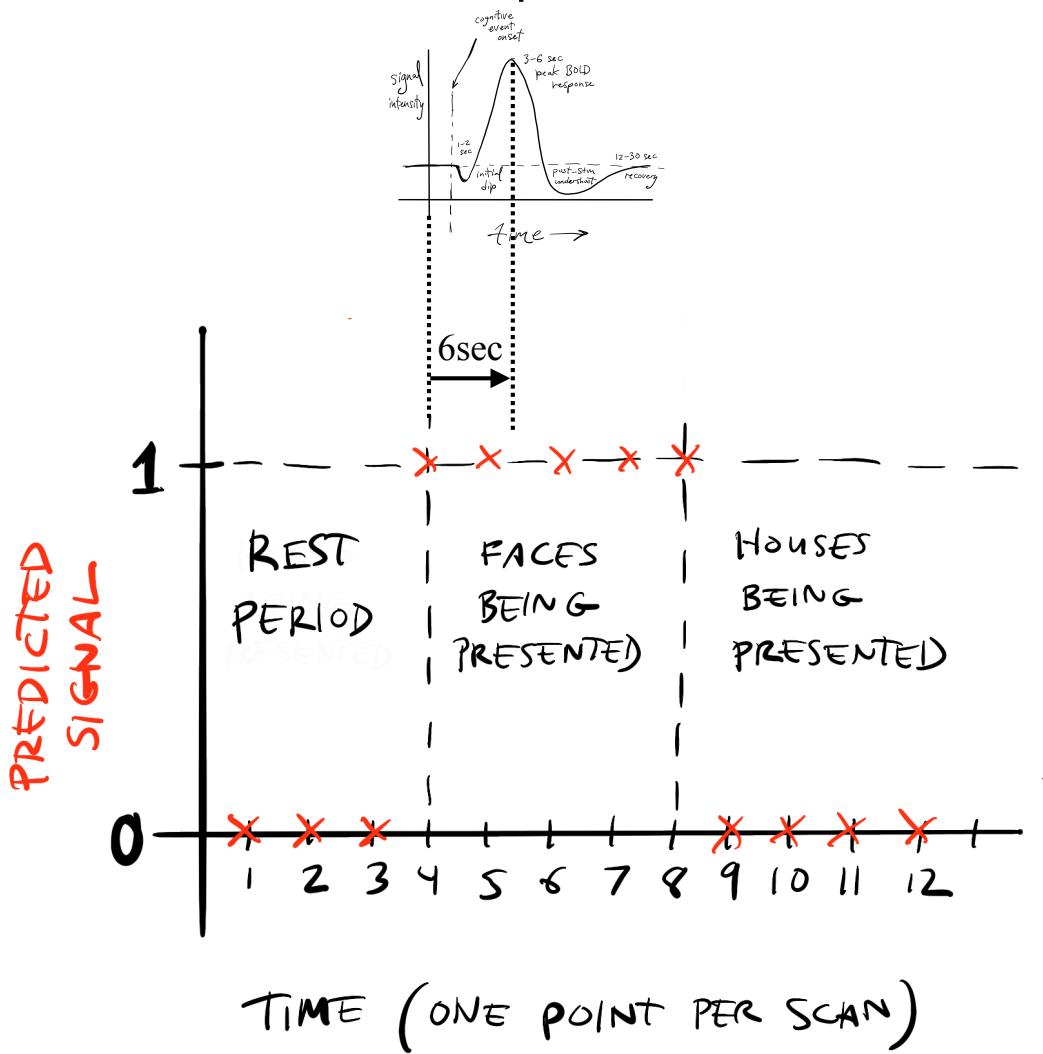
Homework8.ipynb

```
patterns, targets, runs = prep_haxby_data()
```

`prep_haxby_data()`

- opens `haxby_vt_patterns.npy`
- read in `labels.txt`
- "cheap version" of accounting for hemodynamic lag

one approach is to shift predictor function by ~6 sec



Homework8.ipynb

```
patterns, targets, runs = prep_haxby_data()
```

`prep_haxby_data()`

- opens `haxby_vt_patterns.npy`
- read in `labels.txt`
- "cheap version" of accounting for hemodynamic lag
 - each scan was 2.5 sec
 - shift labels by 5 sec (close to 6 sec)
 - which means, shift labels forward by 2 indices

Homework8.ipynb

```
patterns, targets, runs = prep_haxby_data()
```

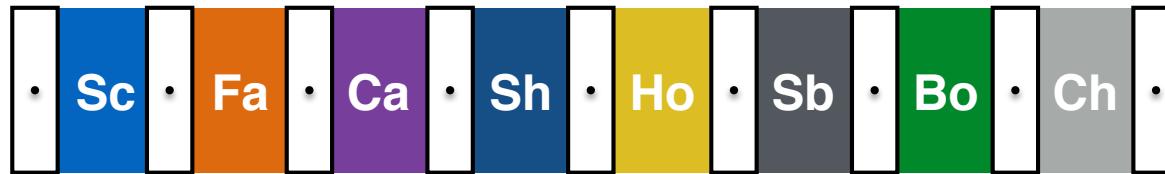
`prep_haxby_data()`

- opens `haxby_vt_patterns.npy`
- read in `labels.txt`
- "cheap version" of accounting for hemodynamic lag
- remove rests from runs

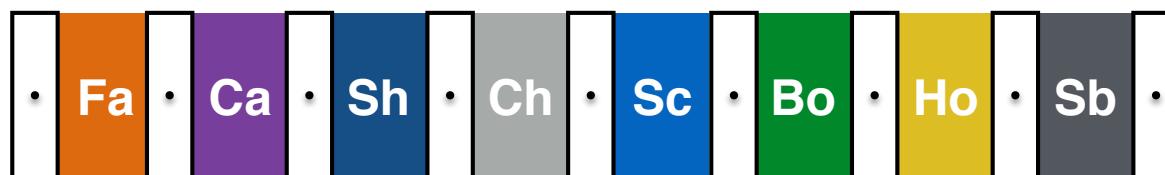
Haxby et al. 2001 details

raw data with rests

run 0

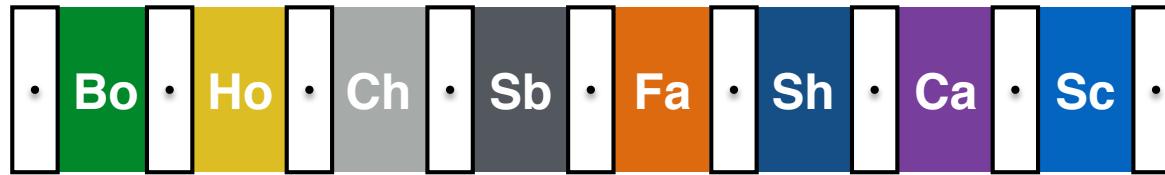


run 1



⋮

run 11



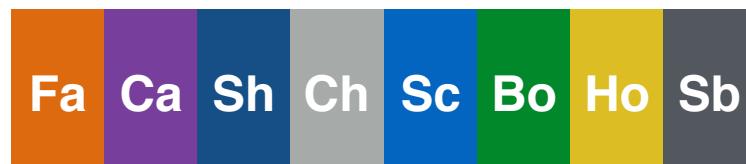
Haxby et al. 2001 details

with rests stripped away *from both labels
and fMRI data*

run 0



run 1



⋮
⋮
⋮

run 11



Homework8.ipynb

```
patterns, targets, runs = prep_haxby_data()
```

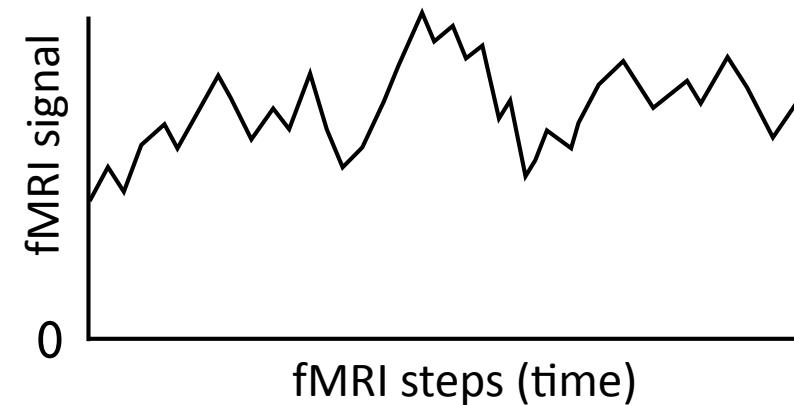
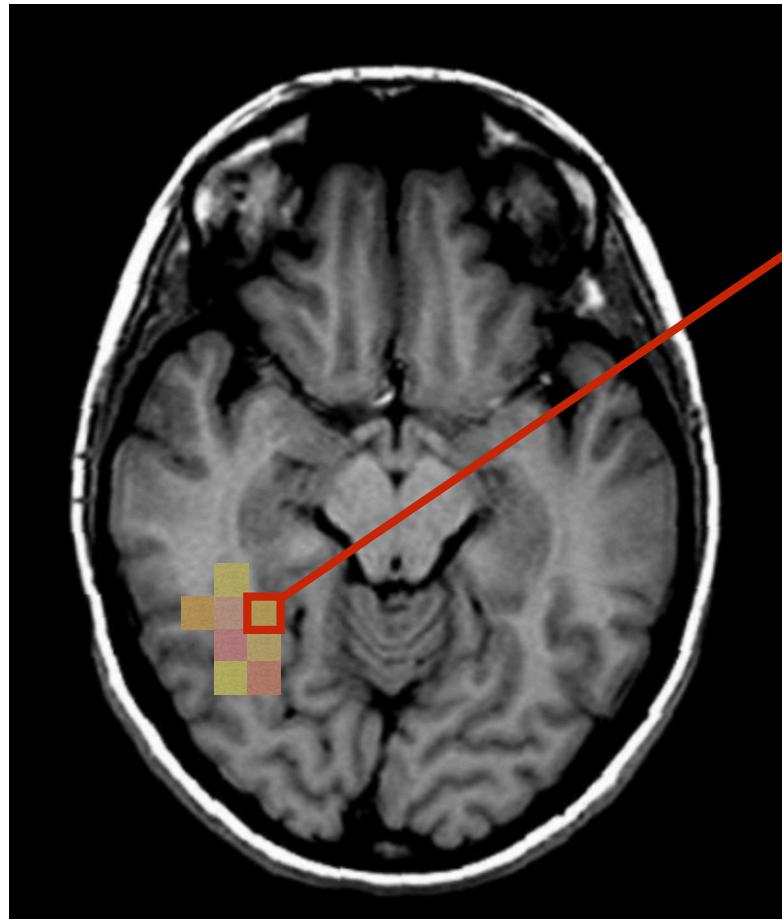
`prep_haxby_data()`

- opens `haxby_vt_patterns.npy`
- read in `labels.txt`
- "cheap version" of accounting for hemodynamic lag
- remove rests from runs
- normalize (z-transform) each voxel

normalize / z-transform

(we've seen before)

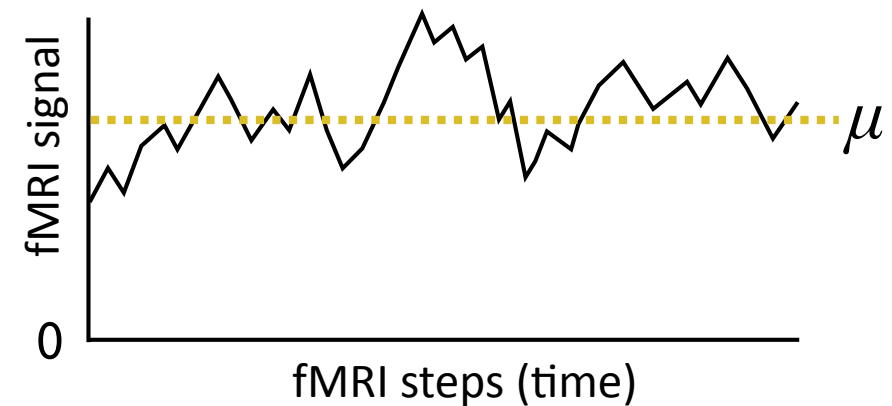
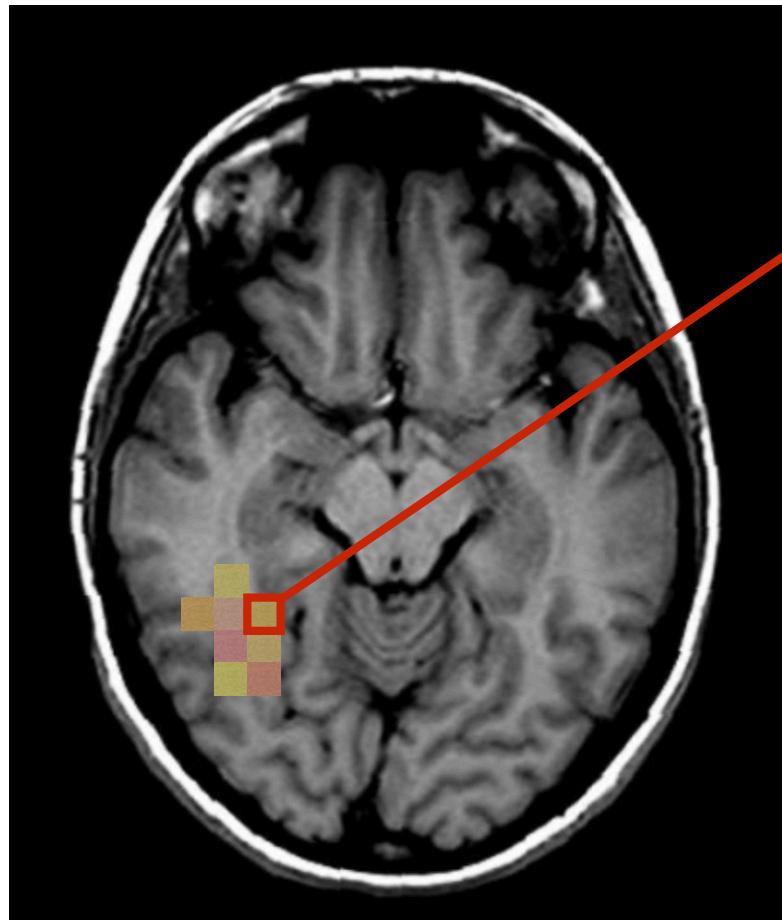
a voxel BEFORE normalization



normalize / z-transform

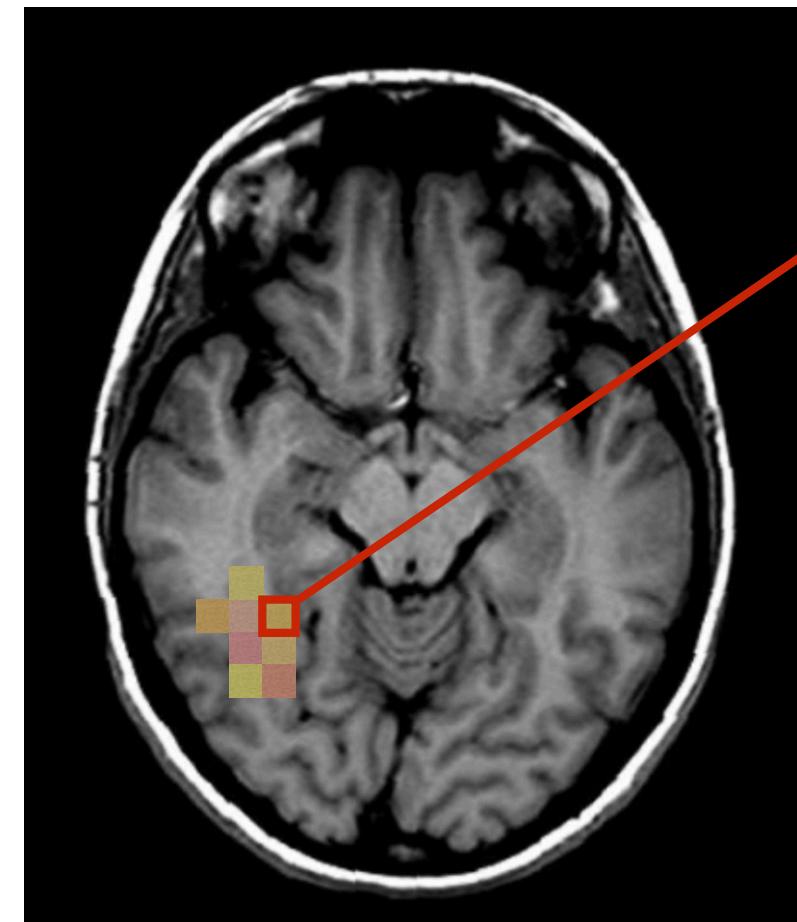
(we've seen before)

a voxel BEFORE normalization

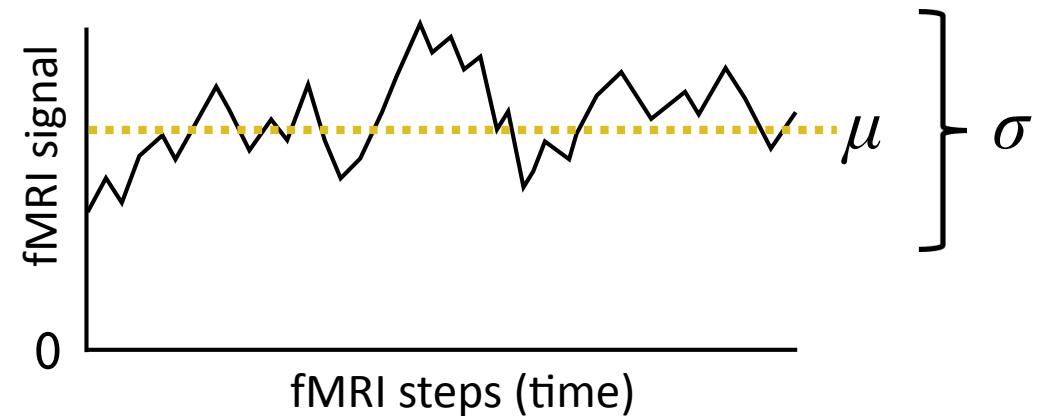


normalize / z-transform

(we've seen before)

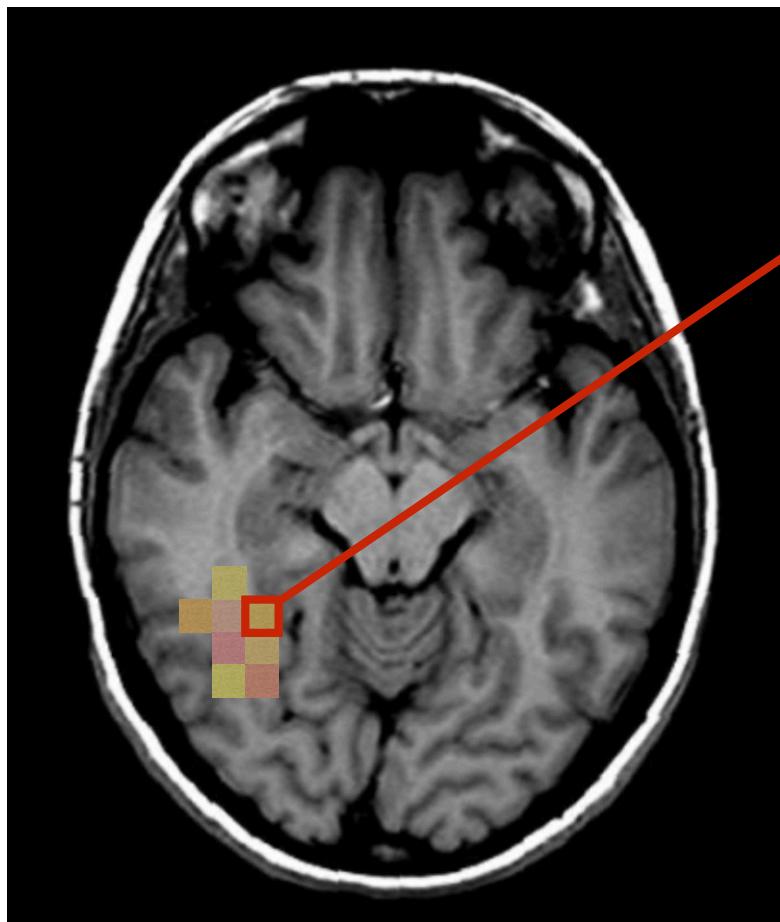


a voxel BEFORE normalization

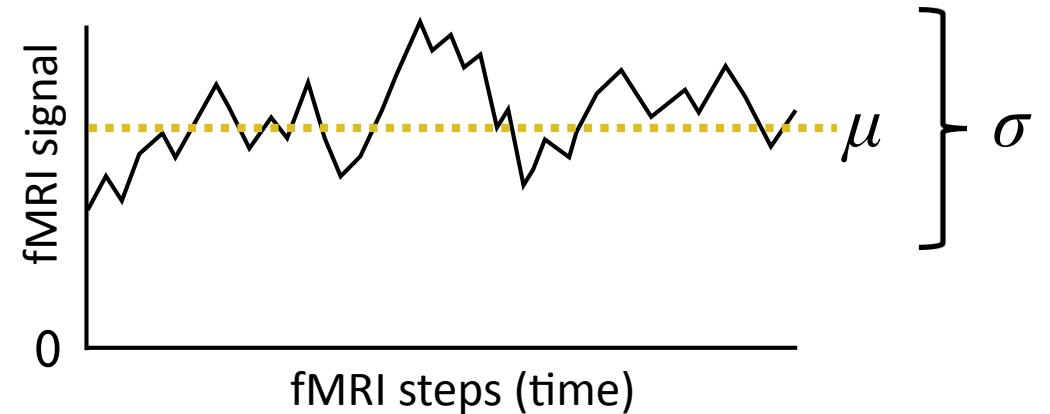


normalize / z-transform

(we've seen before)



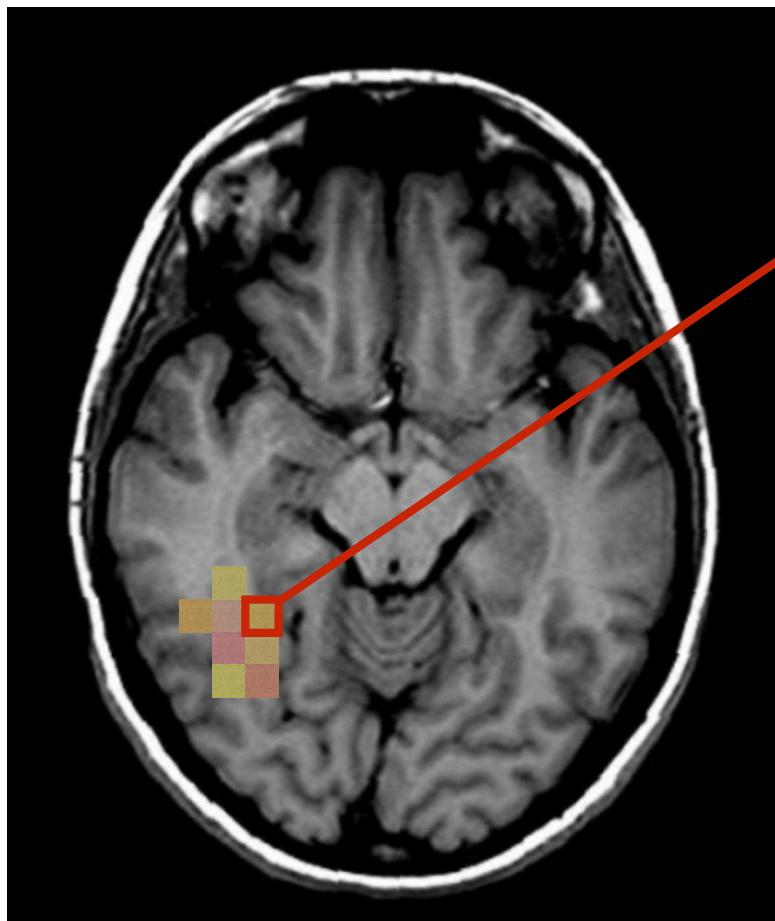
a voxel BEFORE normalization



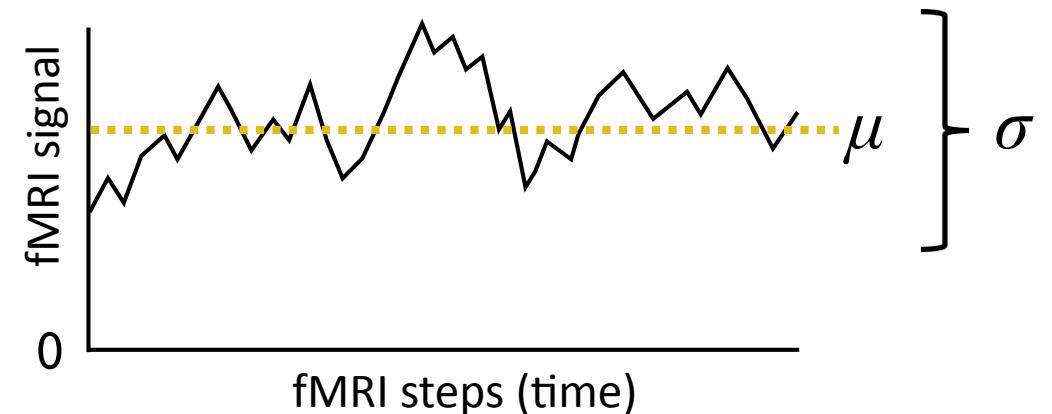
$$x_t = \frac{x_t - \mu}{\sigma}$$

normalize / z-transform

(we've seen before)

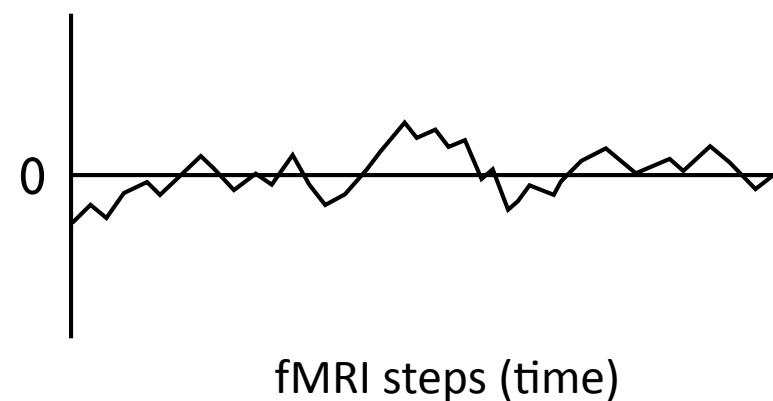


a voxel BEFORE normalization



$$x_t = \frac{x_t - \mu}{\sigma}$$

a voxel AFTER normalization

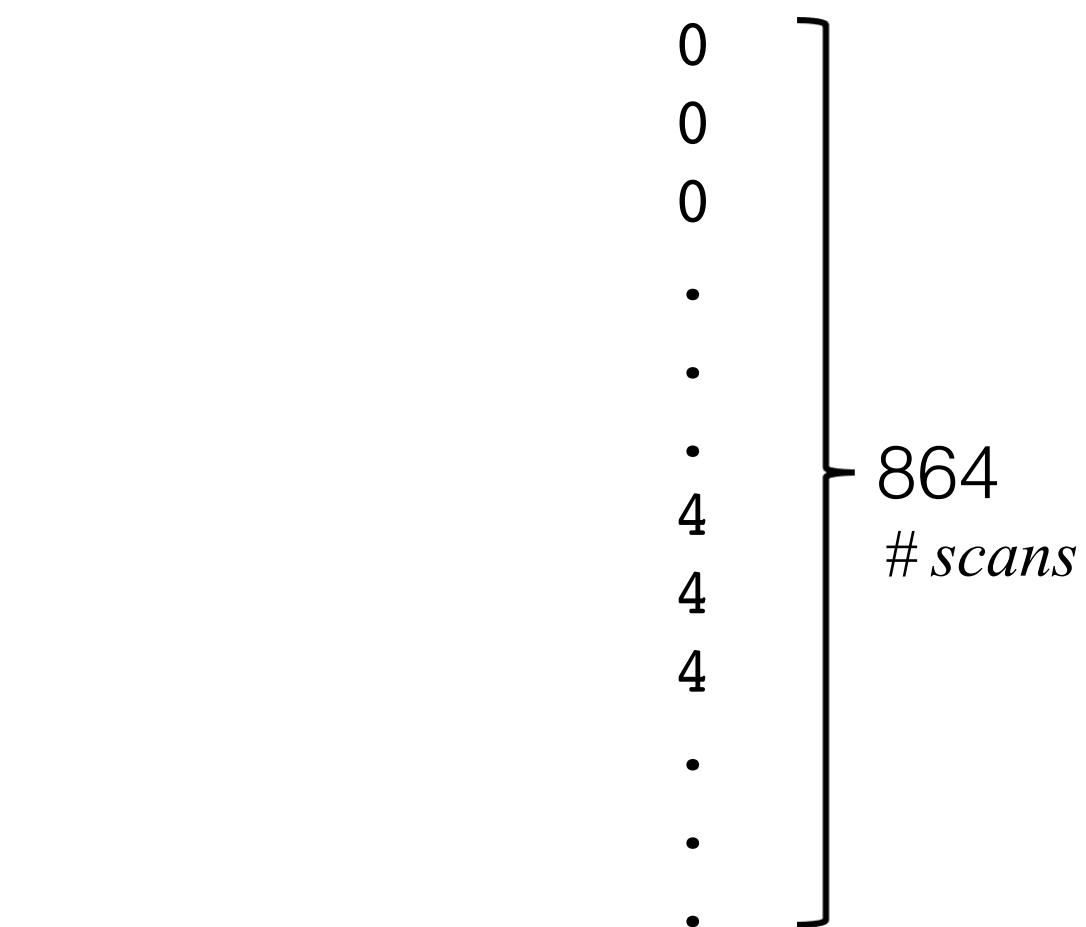


Homework8.ipynb

```
patterns, targets, runs = prep_haxby_data()
```

Homework8.ipynb

```
patterns, targets, runs = prep_haxby_data()
```



Homework8.ipynb

```
patterns, targets, runs_ = prep_haxby_data()
```

categories

8

one hot coding

864
#scans

Homework8.ipynb

```
patterns, targets, runs = prep_haxby_data()
```

voxels

32450

categories

8

1.13	0.23	...
0.03	-0.13	...
0.43	-0.43	...

•

•

•

0.13	-1.23	...
1.13	1.03	...
1.03	-1.02	...

•

•

•

0	0	0	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	0	0	1	0	0	0

•

•

•

0	1	0	0	0	0	0	0
0	1	0	0	0	0	0	0
0	0	0	0	0	0	0	1

0

0

0

•

•

•

4

4

4

•

•

•

864

scans

"feature selection"

we've narrowed things down to voxels in IT cortex (anatomically)

let's narrow to voxels that show some statistically significant modulation with category (eliminating lower-level and non-visual)

```
def feature_selection...
```

```
.
```

```
.
```

```
.
```

```
sps.f_oneway(groups[0], groups[1], groups[2], groups[3],  
             groups[4], groups[5], groups[6], groups[7])
```

simple one-way ANOVA

voxels → # voxels
32450 12432

we'll see where we do this later

classification in MVPA

it looks like we have enough to train a classifier

# voxels			# categories							
32450			8							
1.13	0.23	...	0	0	0	0	1	0	0	0
0.03	-0.13	...	0	0	0	0	1	0	0	0
0.43	-0.43	...	0	0	0	0	1	0	0	0
.			.							
.			.							
.			.							
0.13	-1.23	...	0	1	0	0	0	0	0	0
1.13	1.03	...	0	1	0	0	0	0	0	0
1.03	-1.02	...	0	0	0	0	0	0	1	0
.			.							
.			.							
.			.							

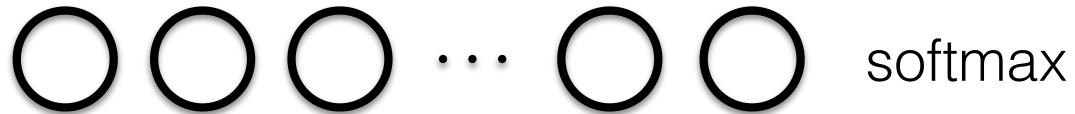
classification in MVPA

it looks like we have enough to train a classifier

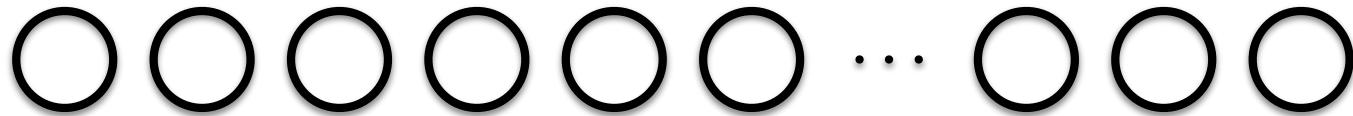
Nin			Nout							
Np	1.13	0.23	...	0	0	0	0	1	0	0
	0.03	-0.13	...	0	0	0	0	1	0	0
	0.43	-0.43	...	0	0	0	0	1	0	0
	.							.		
	.							.		
	.							.		
	0.13	-1.23	...	0	1	0	0	0	0	0
	1.13	1.03	...	0	1	0	0	0	0	0
	1.03	-1.02	...	0	0	0	0	0	0	1
	.							.		
	.							.		
	.							.		

classification in MVPA

categorical cross
entropy loss



densely connected



	Nin			Nout							
Np	1.13	0.23	...	0	0	0	0	1	0	0	0
	0.03	-0.13	...	0	0	0	0	1	0	0	0
	0.43	-0.43	...	0	0	0	0	1	0	0	0
	.			.				.			
	.			.				.			
	.			.				.			
	0.13	-1.23	...	0	1	0	0	0	0	0	0
	1.13	1.03	...	0	1	0	0	0	0	0	0
	1.03	-1.02	...	0	0	0	0	0	0	1	0
	.			.				.			

classification in MVPA

we're not training a classifier to classify some future data ...

we're using a classifier to measure neural activity, to ask whether there are patterns of voxel activity that predict object category

the classifier is a scientific (neural) measurement tool

classification in MVPA

RUN 0	RUN 1	RUN 2	RUN 3
faces	faces	faces	faces
houses	houses	houses	houses
cats	cats	cats	cats
chairs	chairs	chairs	chairs
shoes	shoes	shoes	shoes
bottles	bottles	bottles	bottles
scissors	scissors	scissors	scissors
scrambled	scrambled	scrambled	scrambled

let's keep things simple and assume 4 runs
(rather than the 12 runs we have)

classification in MVPA

RUN 0	RUN 1	RUN 2	RUN 3
faces	faces	faces	faces
houses	houses	houses	houses
cats	cats	cats	cats
chairs	chairs	chairs	chairs
shoes	shoes	shoes	shoes
bottles	bottles	bottles	bottles
scissors	scissors	scissors	scissors
scrambled	scrambled	scrambled	scrambled

null hypothesis is that classifier performs at chance
(no pattern of neural activity predicts object category)

8 categories, chance is $1/8=12.5\%$ classification

classification in MVPA

- training data
- testing data

RUN 0	RUN 1	RUN 2	RUN 3
faces	faces	faces	faces
houses	houses	houses	houses
cats	cats	cats	cats
chairs	chairs	chairs	chairs
shoes	shoes	shoes	shoes
bottles	bottles	bottles	bottles
scissors	scissors	scissors	scissors
scrambled	scrambled	scrambled	scrambled

Haxby essentially used 1/2 the data (odd runs) to compare with other 1/2 of data (even runs) (with a "correlation classifier")

classification in MVPA

- training data
- testing data

	RUN 0	RUN 1	RUN 2	RUN 3
faces				
houses				
cats				
chairs				
shoes				
bottles				
scissors				
scrambled				

we could train a neural network classifier on 1/2 the data
and test on the other 1/2 of the data

classification in MVPA

 training data

 testing data

RUN 0

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 1

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 2

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 3

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

or split up the data this way

classification in MVPA

 training data

 testing data

RUN 0

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 1

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 2

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 3

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

or train on 3/4 and test on a 1/4

which is the best choice?

classification in MVPA

- training data
- testing data

	RUN 0	RUN 1	RUN 2	RUN 3
faces				
houses				
cats				
chairs				
shoes				
bottles				
scissors				
scrambled				

we don't have a lot of data, and the more training data the better, but that means less testing data

classification in MVPA

- training data
- testing data

RUN 0	RUN 1	RUN 2	RUN 3
faces	faces	faces	faces
houses	houses	houses	houses
cats	cats	cats	cats
chairs	chairs	chairs	chairs
shoes	shoes	shoes	shoes
bottles	bottles	bottles	bottles
scissors	scissors	scissors	scissors
scrambled	scrambled	scrambled	scrambled

having more testing data is nice, but that means having less training data for the classifier

Homework8.ipynb

```
patterns, targets, runs = prep_haxby_data()
```

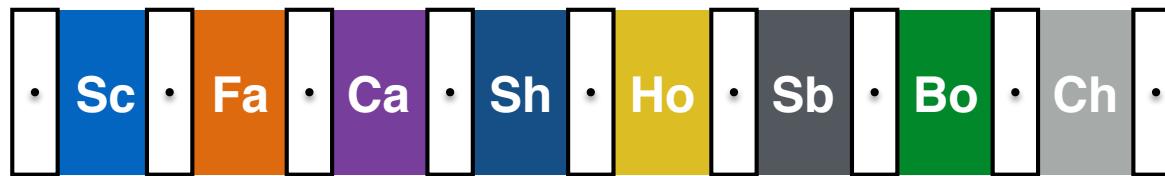
`prep_haxby_data()`

- opens `haxby_vt_patterns.npy`
- read in `labels.txt`
- "cheap version" of accounting for hemodynamic lag
- remove rests from runs

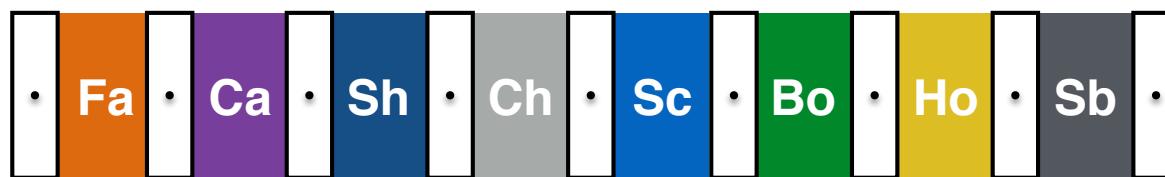
Haxby et al. 2001 details

raw data with rests

run 0

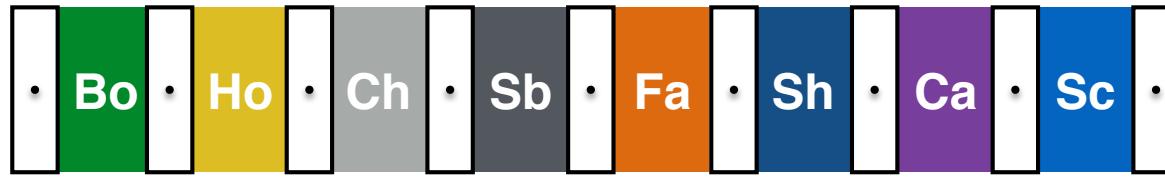


run 1



⋮

run 11



Haxby et al. 2001 details

with rests stripped away *from both labels
and fMRI data*

run 0



run 1



⋮
⋮
⋮

run 11



Homework8.ipynb

```
patterns, targets, runs = prep_haxby_data()
```

voxels

32450

categories

8

1.13	0.23	...
0.03	-0.13	...
0.43	-0.43	...

•

•

•

0.13	-1.23	...
1.13	1.03	...
1.03	-1.02	...

•

•

•

0	0	0	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	0	0	1	0	0	0

•

•

•

0	1	0	0	0	0	0	0
0	1	0	0	0	0	0	0
0	0	0	0	0	0	0	1

0
0
0

4
4
4

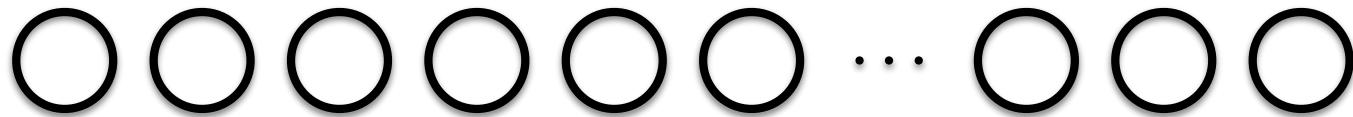
864
scans

classification in MVPA

categorical cross
entropy loss



densely connected



	Nin			Nout							
Np	1.13	0.23	...	0	0	0	0	1	0	0	0
	0.03	-0.13	...	0	0	0	0	1	0	0	0
	0.43	-0.43	...	0	0	0	0	1	0	0	0
	.			.				.			
	.			.				.			
	.			.				.			
	0.13	-1.23	...	0	1	0	0	0	0	0	0
	1.13	1.03	...	0	1	0	0	0	0	0	0
	1.03	-1.02	...	0	0	0	0	0	0	1	0
	.			.				.			

Homework8.ipynb

not much data to constrains lots of parameters

`patterns, targets, runs = prep_haxby_data()`

voxels

32450

categories

8

1.13	0.23	...	0	0	0	0	1	0	0	0	0
0.03	-0.13	...	0	0	0	0	1	0	0	0	0
0.43	-0.43	...	0	0	0	0	1	0	0	0	0
.
0.13	-1.23	...	0	1	0	0	0	0	0	0	4
1.13	1.03	...	0	1	0	0	0	0	0	0	4
1.03	-1.02	...	0	0	0	0	0	0	1	0	4
.

864

scans

classification in MVPA

- training data
- testing data

	RUN 0	RUN 1	RUN 2	RUN 3
faces				
houses				
cats				
chairs				
shoes				
bottles				
scissors				
scrambled				

we don't have a lot of data, and the more training data the better, but that means less testing data

classification in MVPA

- training data
- testing data

RUN 0	RUN 1	RUN 2	RUN 3
faces	faces	faces	faces
houses	houses	houses	houses
cats	cats	cats	cats
chairs	chairs	chairs	chairs
shoes	shoes	shoes	shoes
bottles	bottles	bottles	bottles
scissors	scissors	scissors	scissors
scrambled	scrambled	scrambled	scrambled

having more testing data is nice, but that means having less training data for the classifier

cross-validation

- training data
- testing data

RUN 0	RUN 1	RUN 2	RUN 3
faces	faces	faces	faces
houses	houses	houses	houses
cats	cats	cats	cats
chairs	chairs	chairs	chairs
shoes	shoes	shoes	shoes
bottles	bottles	bottles	bottles
scissors	scissors	scissors	scissors
scrambled	scrambled	scrambled	scrambled

each iteration of cross-validation is called a "fold"

with 4 runs, we'd have 4 folds
(with 12 runs, we have 12 folds)

cross-validation

41%

RUN 0	RUN 1	RUN 2	RUN 3
faces	faces	faces	faces
houses	houses	houses	houses
cats	cats	cats	cats
chairs	chairs	chairs	chairs
shoes	shoes	shoes	shoes
bottles	bottles	bottles	bottles
scissors	scissors	scissors	scissors
scrambled	scrambled	scrambled	scrambled

first fold: train on runs 0-3, test on run 4

record classification accuracy for this fold

cross-validation

38%

41%

RUN 0	RUN 1	RUN 2	RUN 3
faces	faces	faces	faces
houses	houses	houses	houses
cats	cats	cats	cats
chairs	chairs	chairs	chairs
shoes	shoes	shoes	shoes
bottles	bottles	bottles	bottles
scissors	scissors	scissors	scissors
scrambled	scrambled	scrambled	scrambled

second fold

cross-validation

	45%	38%	41%				
RUN 0	faces houses cats chairs shoes bottles scissors scrambled	RUN 1	faces houses cats chairs shoes bottles scissors scrambled	RUN 2	faces houses cats chairs shoes bottles scissors scrambled	RUN 3	faces houses cats chairs shoes bottles scissors scrambled

third fold

cross-validation

42%

45%

38%

41%

RUN 0

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 1

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 2

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 3

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

fourth fold

cross-validation

42%

45%

38%

41%

RUN 0

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 1

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 2

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 3

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

overall performance

41.5%

(remember, chance is 12.5%)

cross-validation

overall performance

41.5%

RUN 0

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 1

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 2

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 3

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

cross-validation essentially lets you use all of
the data to both train and test classifiers

cross-validation with fMRI

RUN 0

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 1

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 2

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

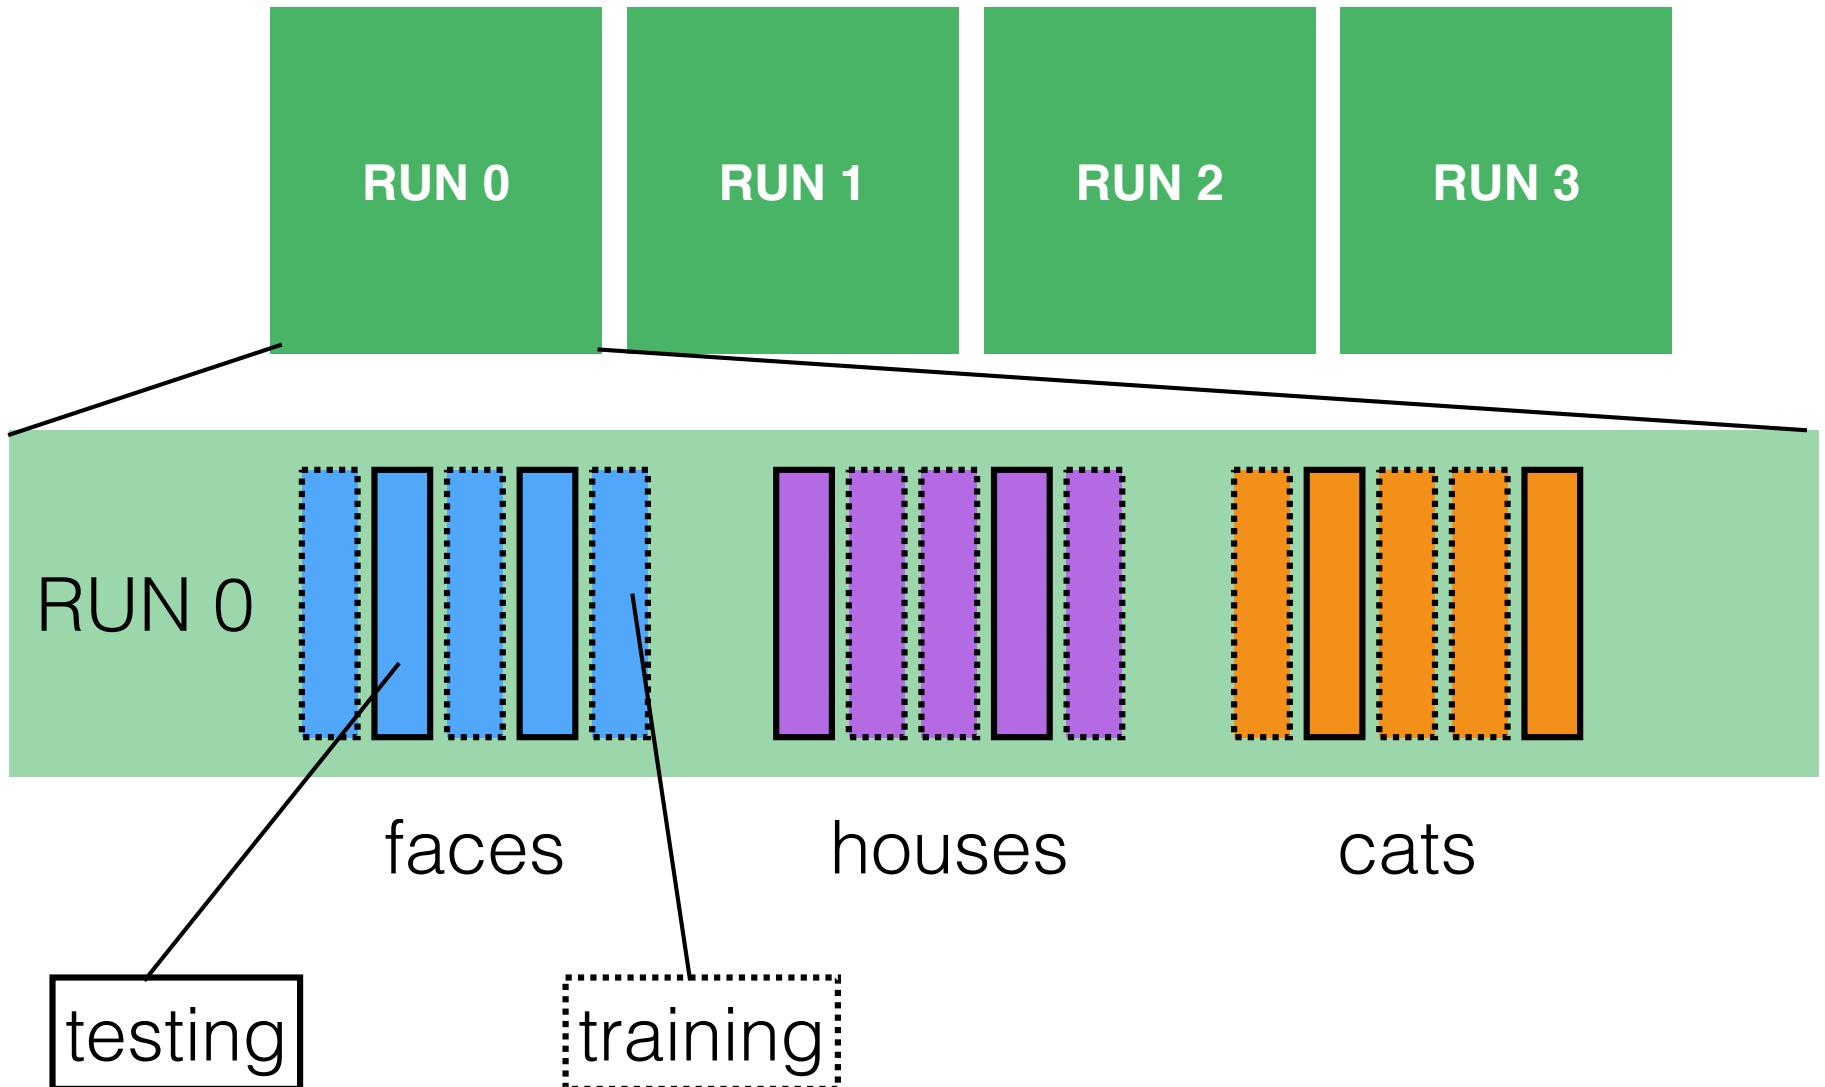
RUN 3

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

note that here we're selecting folds based on complete runs
(which is a recommended method for fMRI)

cross-validation with fMRI

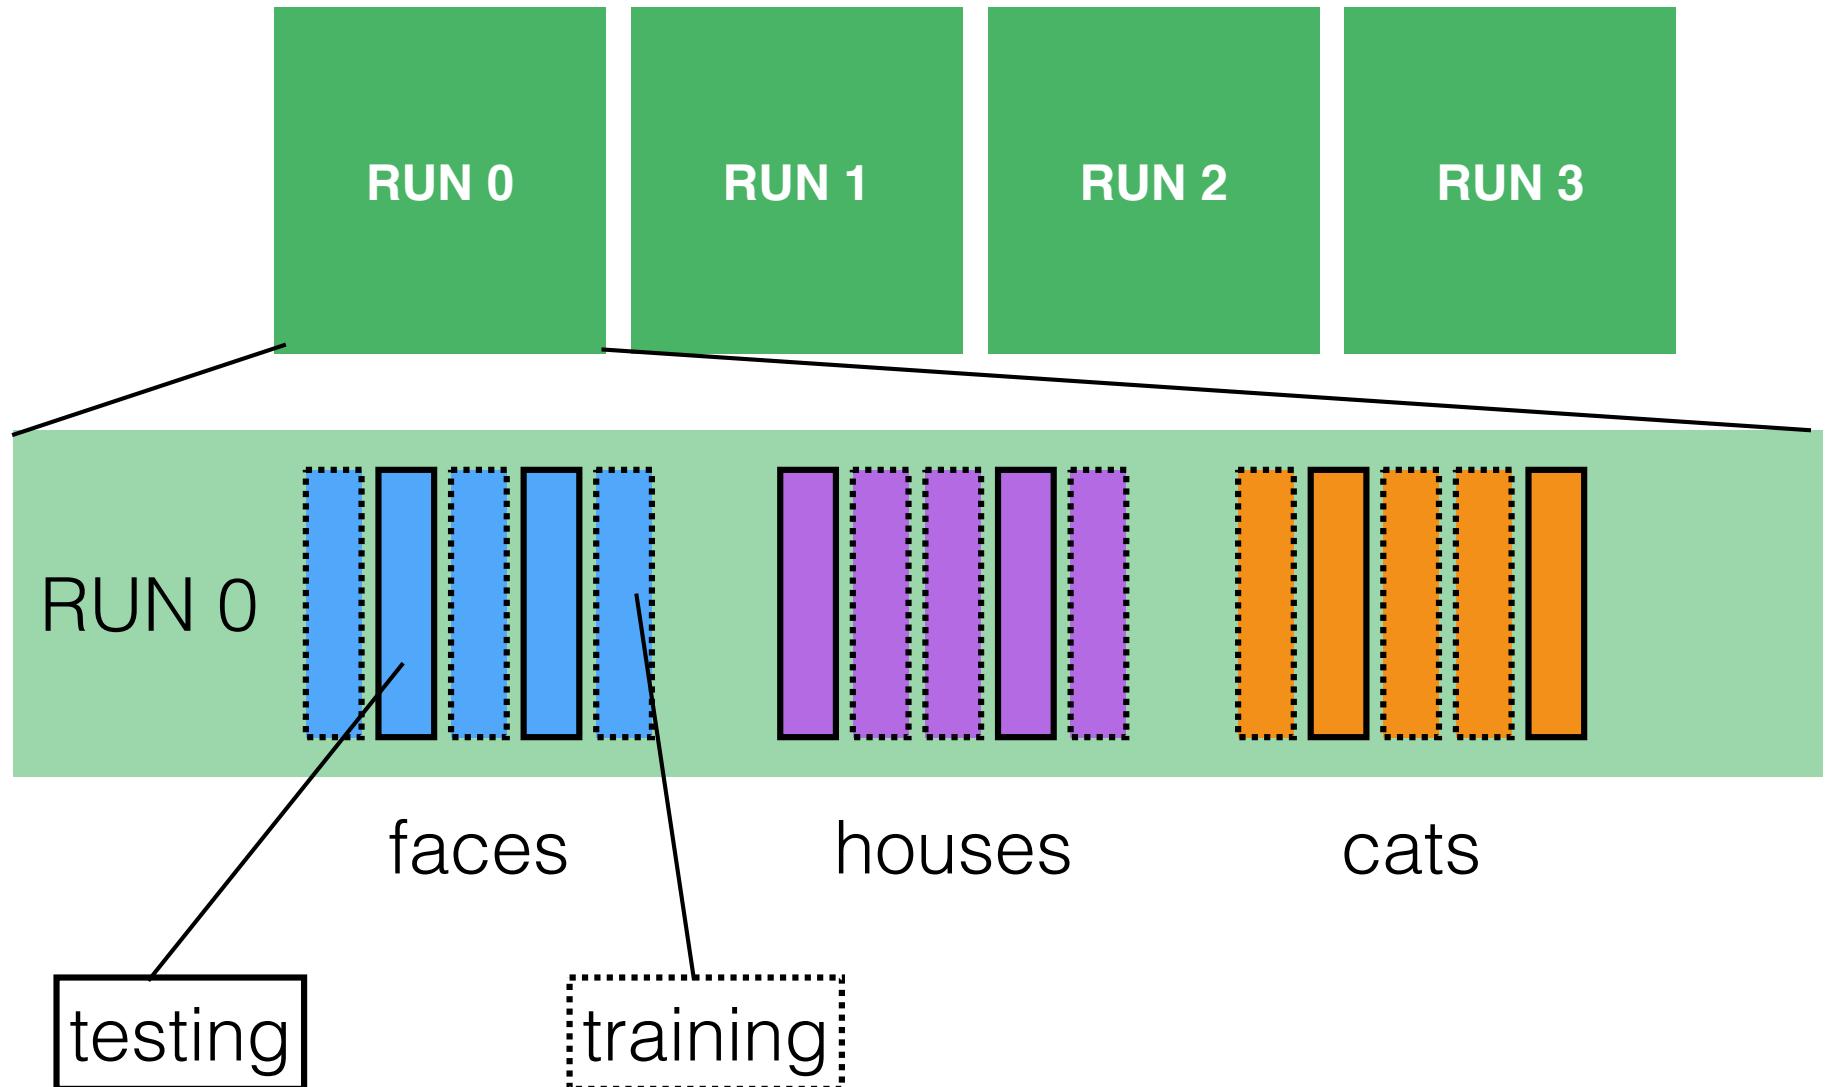
why not select folds within a run?



cross-validation with fMRI

(often from physical, non-physiological sources)

there is substantial "temporal autocorrelation" in fMRI data,
which could spuriously inflate classifier performance



Homework 8 (PART ONE)

implement a 12-fold cross-validation of **MVPA** using a **softmax classifier**

do patterns of voxel activity predict object category

single-layer neural network with softmax and categorical cross entropy

set up data structures to store
classifier performance

set up data structures to store
classifier performance

```
for r in range(n_runs):
```

loop over the 12 folds

```
set up data structures to store  
classifier performance
```

```
for r in range(n_runs):
```

```
    set up the fold r
```

```
set up data structures to store  
classifier performance
```

```
for r in range(n_runs):
```

```
    set up the fold r
```

```
        do feature selection for this fold
```

```
set up data structures to store  
classifier performance
```

```
for r in range(n_runs):
```

```
    set up the fold r
```

```
        do feature selection for this fold
```

```
        train classifier for this fold
```

```
set up data structures to store  
classifier performance
```

```
for r in range(n_runs):
```

```
    set up the fold r
```

```
    do feature selection for this fold
```

```
    train classifier for this fold
```

```
    test classifier for this fold
```

```
set up data structures to store  
classifier performance
```

```
for r in range(n_runs):
```

```
    set up the fold r
```

```
    do feature selection for this fold
```

```
    train classifier for this fold
```

```
    test classifier for this fold
```

```
    save classifier performance
```

```
set up data structures to store  
classifier performance
```

```
for r in range(n_runs):
```

```
    set up the fold r
```

```
    do feature selection for this fold
```

```
    train classifier for this fold
```

```
    test classifier for this fold
```

```
    save classifier performance
```

```
report average of classifiers
```

set up data structures to store
classifier performance

```
for i in range(n_iters):  
    for r in range(n_runs):
```

*do this many times
(since neural networks stochastic)*

set up the fold r

do feature selection for this fold

train classifier for this fold

test classifier for this fold

save classifier performance

report average of classifiers

set up data structures to store
classifier performance

```
for i in range(n_iters):      in practice, n_iters = 100  
                            for Homework, n_iters = 5
```

```
    for r in range(n_runs):
```

set up the fold r

do feature selection for this fold

train classifier for this fold

test classifier for this fold

save classifier performance

report average of classifiers

```
set up data structures to store  
classifier performance
```

```
for i in range(n_iters):  
    for r in range(n_runs):
```

```
        set up the fold r
```

```
        do feature selection for this fold
```

```
        train classifier for this fold
```

```
        test classifier for this fold
```

```
        save classifier performance
```

```
    report average of classifiers
```

**make each of these
a separate function**

set up the fold r



training data



testing data

RUN 0

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 1

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

RUN 2

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

...

RUN 11

faces
houses
cats
chairs
shoes
bottles
scissors
scrambled

for example, suppose we're leaving out run 1

set up the fold r

see **Homework8.ipynb**

set up the fold r

```
# an example "leaving out" a particular run

example_test_run = 1

train_these = runs != example_test_run
test_these = runs == example_test_run

# an example of logical indexing
train_patterns = patterns[train_these, :]
train_targets = targets[train_these, :]

test_patterns = patterns[test_these, :]
test_targets = targets[test_these, :]
```

do feature selection for this fold

```
fs_train_patterns, fs_test_patterns =  
    feature_selection(train_patterns,  
                      train_targets,  
                      test_patterns)
```



*these will have different numbers
of voxels depending on run*

train classifier for this fold

- create classifier in Keras
 - softmax activation function
 - categorical cross entropy for loss function

train classifier for this fold

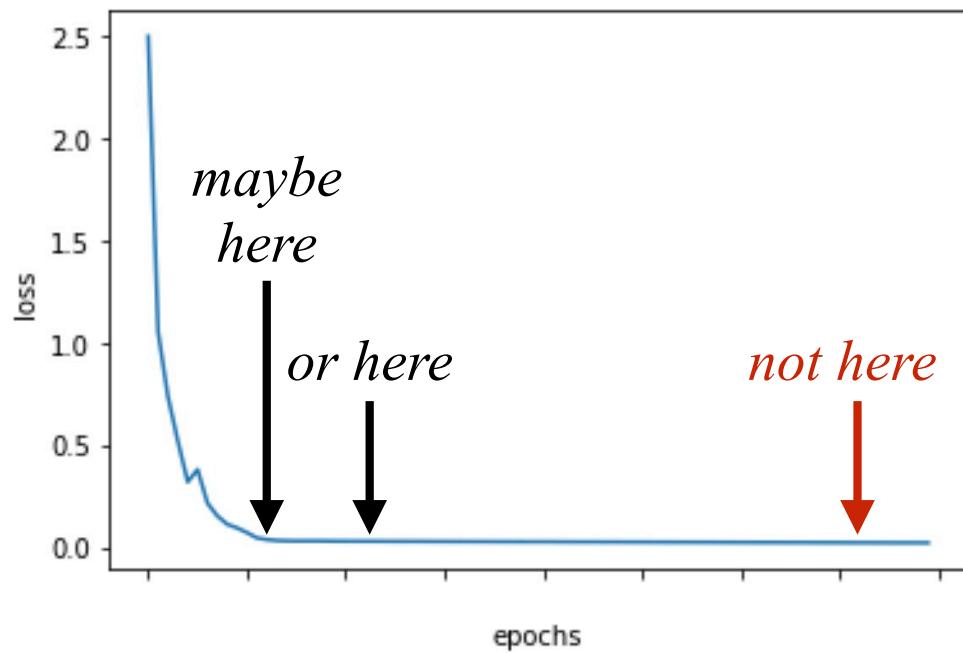
- create classifier in Keras
 - softmax activation function
 - categorical cross entropy for loss function
 - use no validation patterns
 - *too little data to set aside validation patterns*
 - *not creating a classification "engine"*
 - *cross-validation fills role of evaluating*

train classifier for this fold

- create classifier in Keras
 - softmax activation function
 - categorical cross entropy for loss function
 - use no validation patterns
 - use L2 regularization to help avoid overfitting
 - *when you add a layer, include:*
 - `kernel_regularizer=tf.keras.regularizers.l2(0.001)`

train classifier for this fold

- create classifier in Keras
 - softmax activation function
 - categorical cross entropy for loss function
 - use no validation patterns
 - use L2 regularization to help avoid overfitting
- train classifier
 - pick a reasonable number of training epochs



train classifier for this fold

- create classifier in Keras
 - softmax activation function
 - categorical cross entropy for loss function
 - use no validation patterns
 - use L2 regularization to help avoid overfitting
- train classifier
 - pick a reasonable number of training epochs
 - train with `fs_train_patterns` and `train_targets` (see `Homework8.ipynb`)

test classifier for this fold

- test classifier
 - using `network.evaluate()` method for overall performance
 - using `network.predict()` method to get classification of individual test items
 - test with `fs_test_patterns` and `test_targets` (see `Homework8.ipynb`)

save classifier performance

- store in whatever data structure you've set up

report average of classifiers

- average across the 12 folds

testing for statistical significant?

remember, chance performance (1/8) is 12.5%

is 20% classification accuracy greater than chance?

what about 50%

what about 16%

in practice, you could have, for example, chance at 33% and
see *statistically significant* classification at 36%

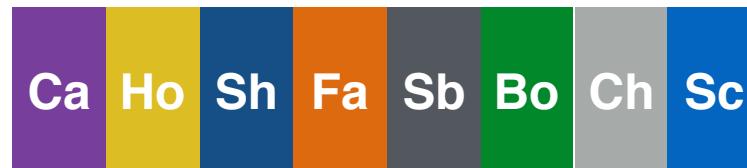
with "data science methods" like this, you often cannot
rely on some standard (parametric) statistical test

permutation test

we're asking if we can predict the object category (label) from patterns of fMRI voxel activity

of course, when doing that, the label (chair, face, etc.) is consistently associated with measures of brain activity (fMRI) when the person is looking at that kind of object

run 3



permutation creates what you might expect "by chance"

permutation test

run 0



remember that each of these contains category labels and fMRI data

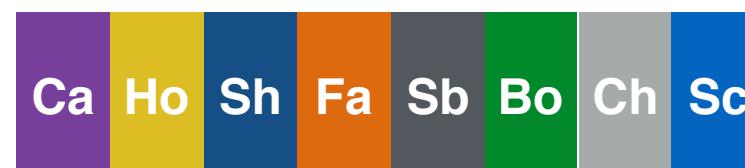
run 1



run 2



run 3



⋮

run 11



permutation test

what if we were to scramble the labels associated with different brain scans (don't touch brain scans)?

e.g., on one run, voxel measures when a face is shown would be labeled as being trials when a chair was shown, and on another run, voxel measures when a face is shown would be labeled as being trials when a shoe was shown

the performance for a classifier trained on that random shuffling of labels would indicate the level of performance you would expect for a classifier "by chance"

permutation test

assume the colors represent the fMRI scans and letter represent labels

run 0

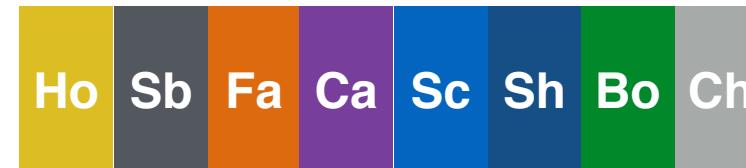


permute the labels

run 1



run 2



run 3



⋮

run 11



permutation test

run 0



run 1



permute the labels

run 2



run 3



⋮

run 11



permutation test

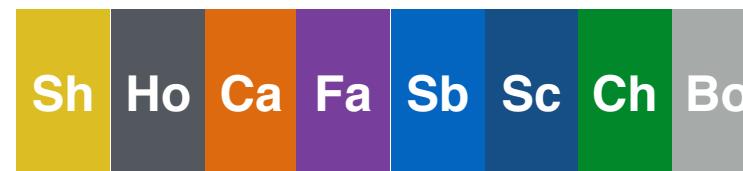
run 0



run 1



run 2



permute the labels

run 3



⋮

run 11



permutation test

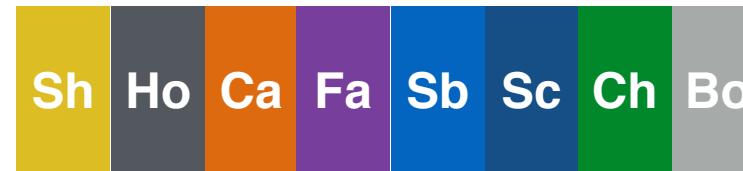
run 0



run 1



run 2

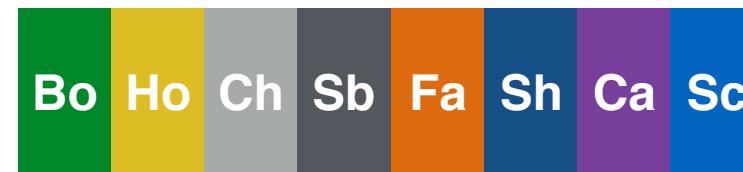


run 3



⋮
⋮
⋮

run 11



permute the labels

permutation test

run 0



run 1



run 2



run 3



⋮
⋮
⋮

run 11



permute the labels

permutation test

```
for p in range(n_perms):
```

permute labels

*in practice, n_perms = 5000
for Homework, n_perms = 10
(but should work for any values of n_perms)*

set up data structures to store
classifier performance

```
for i in range(n_iters):  
    for r in range(n_runs):
```

set up the fold r

do feature selection for this fold

train linear classifier for this fold

test linear classifier for this fold

save classifier performance

report average of classifiers

save average performance to
array to calculate chance range

permutation test

```
for p in range(n_perms):  
    permute labels  
  
    set up data structures to store  
    classifier performance  
  
    for i in range(n_iters):  
        for r in range(n_runs):  
            set up the fold r  
            do feature selection for this fold  
            train linear classifier for this fold  
            test linear classifier for this fold  
            save classifier performance  
  
    report average of classifiers  
  
    save average performance to  
    array to calculate chance range
```

*gave you a function to do this
in Homework8.ipynb*

permutation test

```
for p in range(n_perms):  
    permute labels  
  
    set up data structures to store  
    classifier performance  
  
    for i in range(n_iters):  
        for r in range(n_runs):  
            set up the fold r  
            do feature selection for this fold  
            train linear classifier for this fold  
            test linear classifier for this fold  
            save classifier performance  
  
    report average of classifiers  
  
    save average performance to  
    array to calculate chance range
```

*need to do this part with the
permuted labels*

permutation test

```
for p in range(n_perms):  
    permute labels  
  
    set up data structures to store  
    classifier performance  
  
    for i in range(n_iters):  
        for r in range(n_runs):  
            set up the fold r  
            do feature selection for this fold  
            train linear classifier for this fold  
            test linear classifier for this fold  
            save classifier performance  
  
    report average of classifiers  
  
    save average performance to  
    array to calculate chance range
```



*array of n_perms values
of classifier performance for
permuted labels*

permutation test

with true labels

classification accuracy = 45%

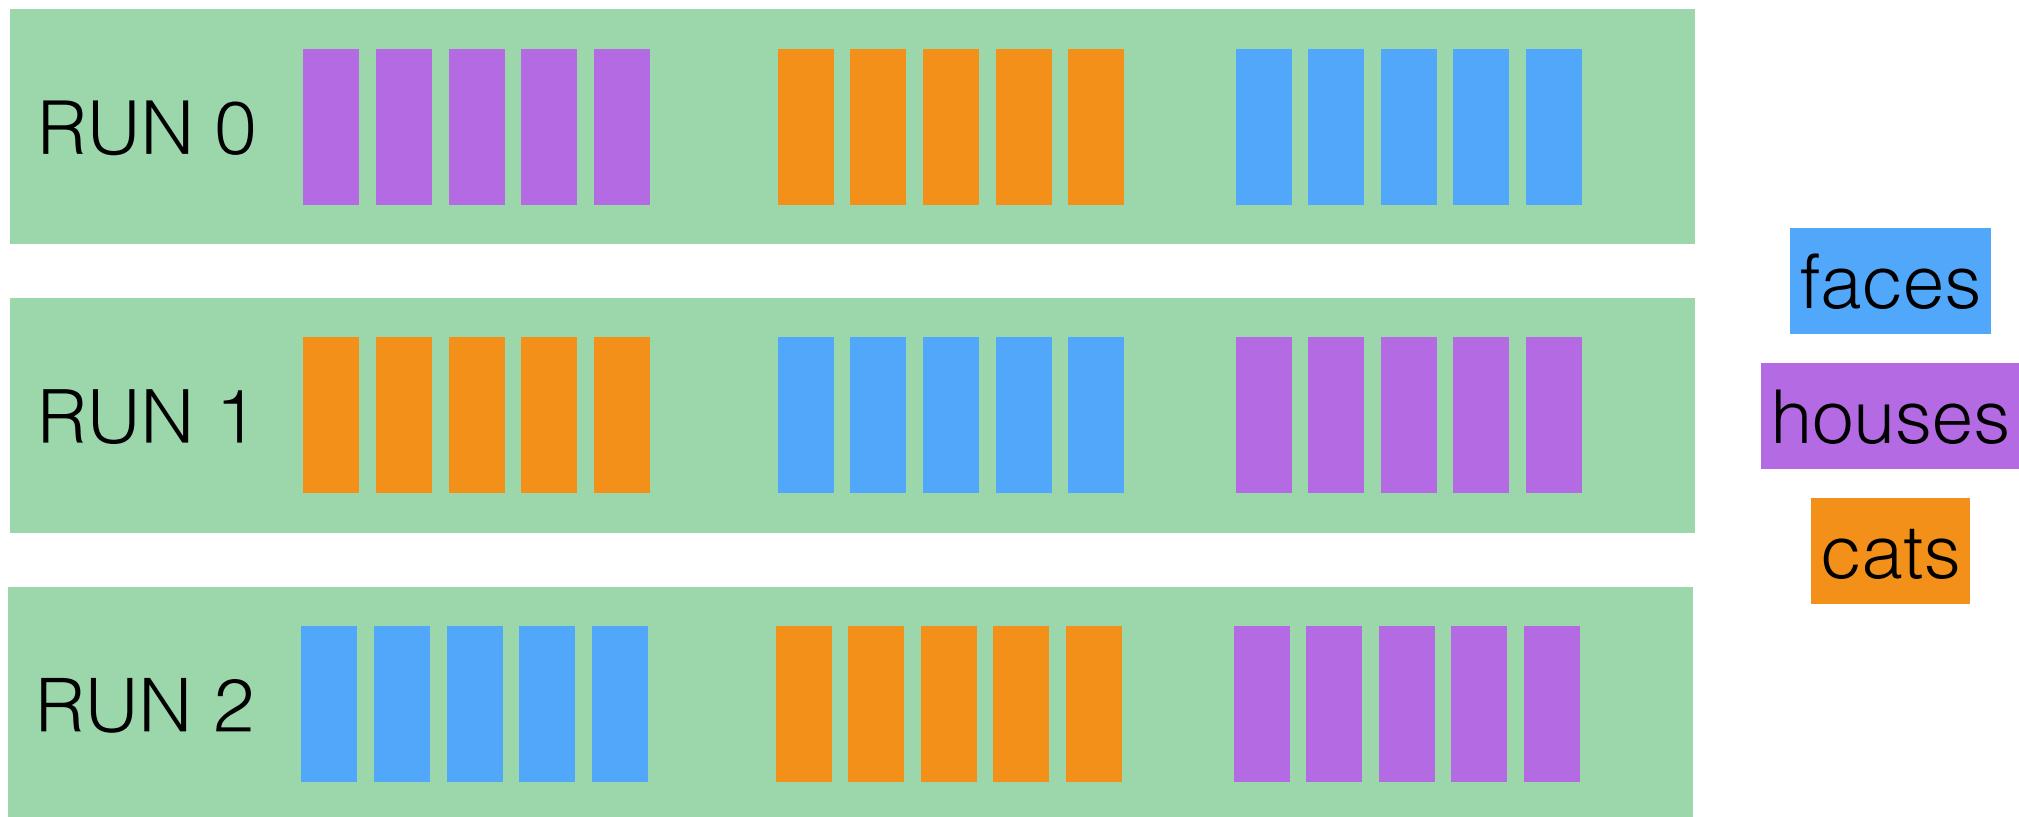


average performance across all runs = “observed accuracy”

permutation test

with scrambled labels

classification accuracy = 14%

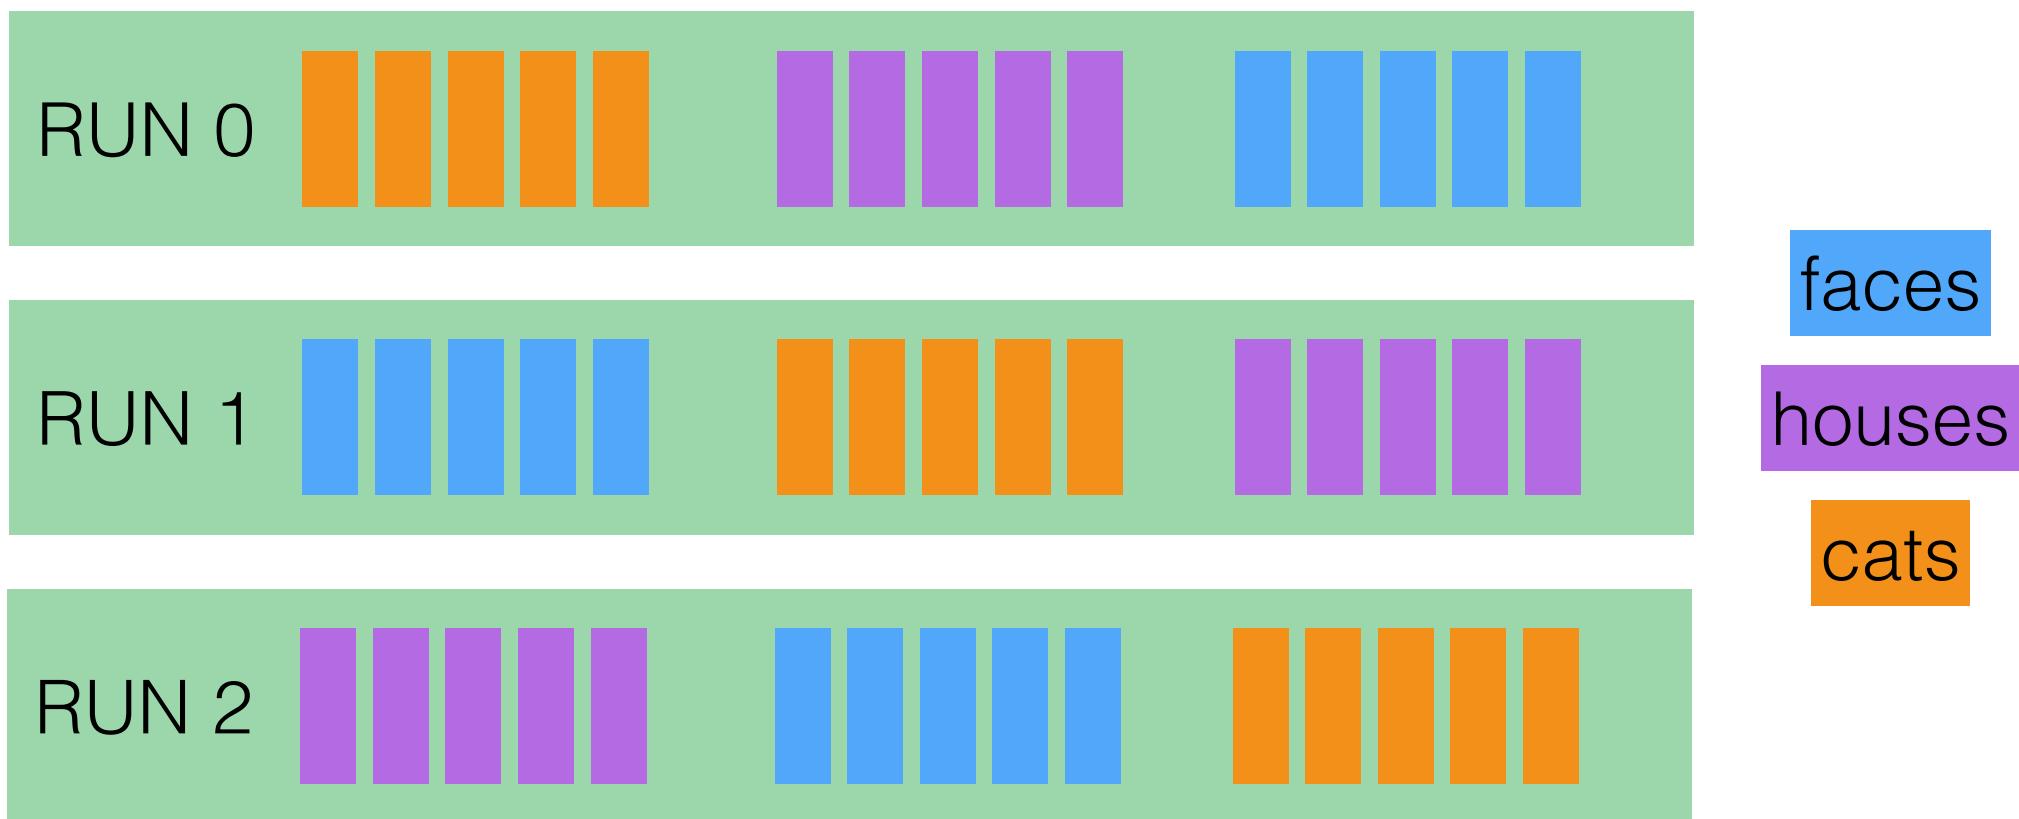


average performance across all runs = “scrambled accuracy #1”

permutation test

with scrambled labels

classification accuracy = 12%



average performance across all runs = “scrambled accuracy #2”

permutation test

observed accuracy = 45%

scrambled accuracy = [14.0,



scrambled accuracy #1

permutation test

observed accuracy = 45%

scrambled accuracy = [14.0, 12.0,



scrambled accuracy #2

permutation test

observed accuracy = 45%

scrambled accuracy = [14.0, 12.0, 8.5,



scrambled accuracy #2

permutation test

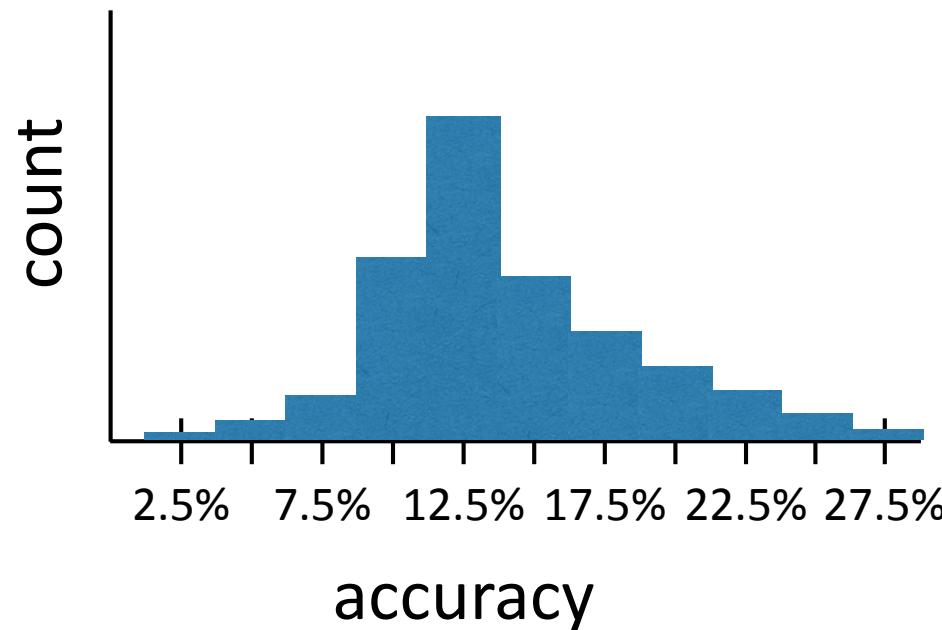
observed accuracy = 45%

scrambled accuracy = [14.0, 12.0, 8.5, 14.3, 8.3, 17.4, 12.9, 11.1, ...]

permutation test

observed accuracy = 45%

scrambled accuracy = [14.0, 12.0, 8.5, 14.3, 8.3, 17.4, 12.9, 11.1, ...]

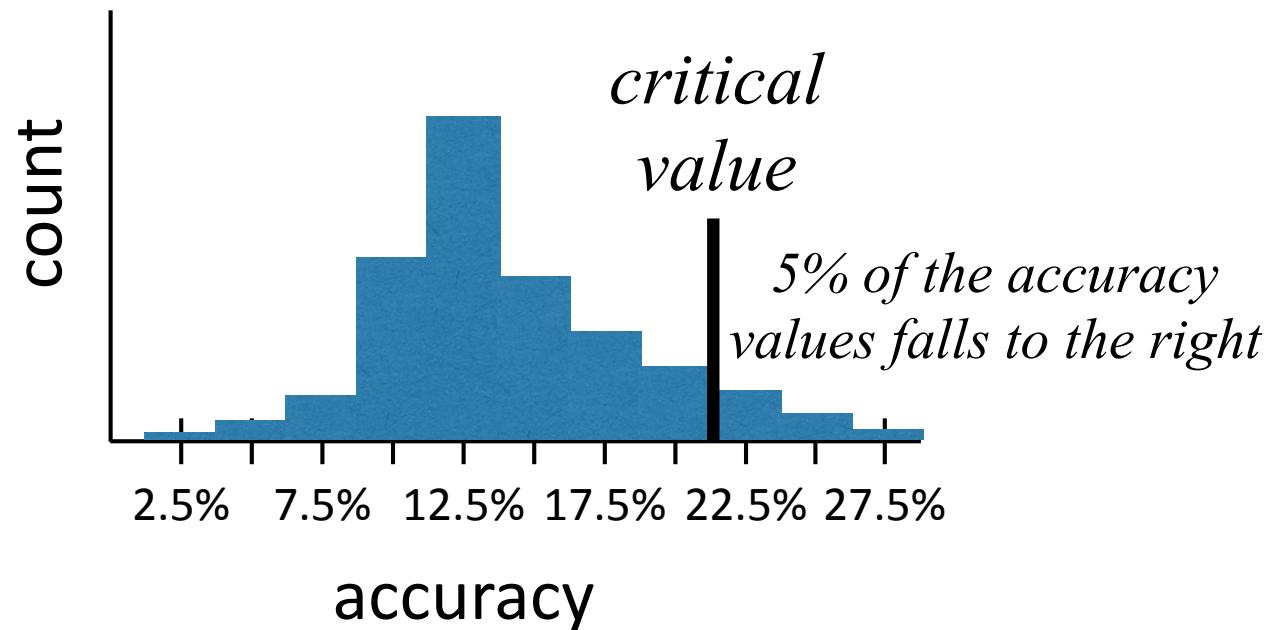


histogram of "chance" performance across
1000 permuted replications of the data

permutation test

observed accuracy = 45%

scrambled accuracy = [14.0, 12.0, 8.5, 14.3, 8.3, 17.4, 12.9, 11.1, ...]

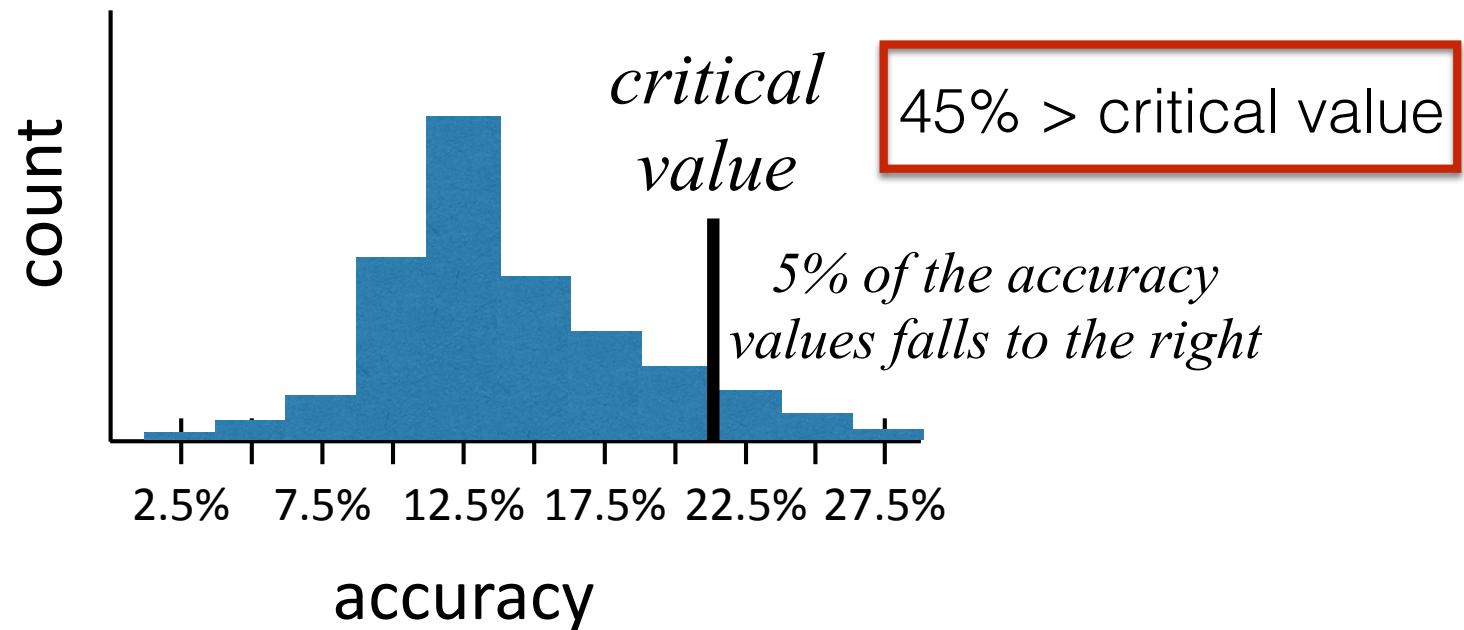


histogram of "chance" performance across
1000 permuted replications of the data

permutation test

observed accuracy = 45%

scrambled accuracy = [14.0, 12.0, 8.5, 14.3, 8.3, 17.4, 12.9, 11.1, ...]



histogram of "chance" performance across
1000 permuted replications of the data