


on Brightspace

Required Video (link on Brightspace)

(for today) *But What *is* a Neural Network?*

by 3blue1brown, Grant Sanderson




Homework 2 (on Brightspace, in today's slides)

Due Tue Sep 14

Simple Models of Neurons

Brief Overview of Neuron Structure and Function


Neuron




Dendrites **Cell Body**

Axon


Synapses




(B) Retinal bipolar cell




(C) Retinal ganglion cell




(D) Retinal amacrine cell




(E) Cortical pyramidal cell




(F) Cerebellar Purkinje cells





functional similarities despite morphological differences





we will not discuss different types of neurons,
but instead model idealized neurons






neurotransmitters cause certain channels to open, allowing charged ions to flow at certain times

ions flow, which changes the voltage of the neuron




EPSPs (excitatory postsynaptic potentials) positive changes in membrane potential caused by excitatory pre-synaptic neurons (e.g., glutamate and NMDA receptors)

IPSPs (inhibitory postsynaptic potentials) negative changes in membrane potential caused by inhibitory pre-synaptic neurons (e.g., GABA and GABA receptors)




inputs can be
positive (+) or
negative(-)



EPSPs (excitatory postsynaptic potentials)
positive changes in membrane potential
caused by excitatory pre-synaptic neurons
(e.g., glutamate and NMDA receptors)


IPSPs (inhibitory postsynaptic potentials)
negative changes in membrane potential
caused by inhibitory pre-synaptic neurons
(e.g., GABA and GABA receptors)



**spatio-temporal summation
of many EPSPs and IPSPs**

from many of synaptic stimulations
in close spatio-temporal proximity

potentials spread, but dissipate
over space and time because
neurons not good conductors




**inputs integrate
at the cell body
- they are added
together**

**spatio-temporal summation
of many EPSPs and IPSPs**


from many of synaptic stimulations
in close spatio-temporal proximity

potentials spread, but dissipate
over space and time because
neurons not good conductors

(A)



(B)




1 EPSP


2 EPSPs

1 IPSP


EPSPs + IPSP

(B)





**an action potential is generated
if net input is greater than
an activation threshold**

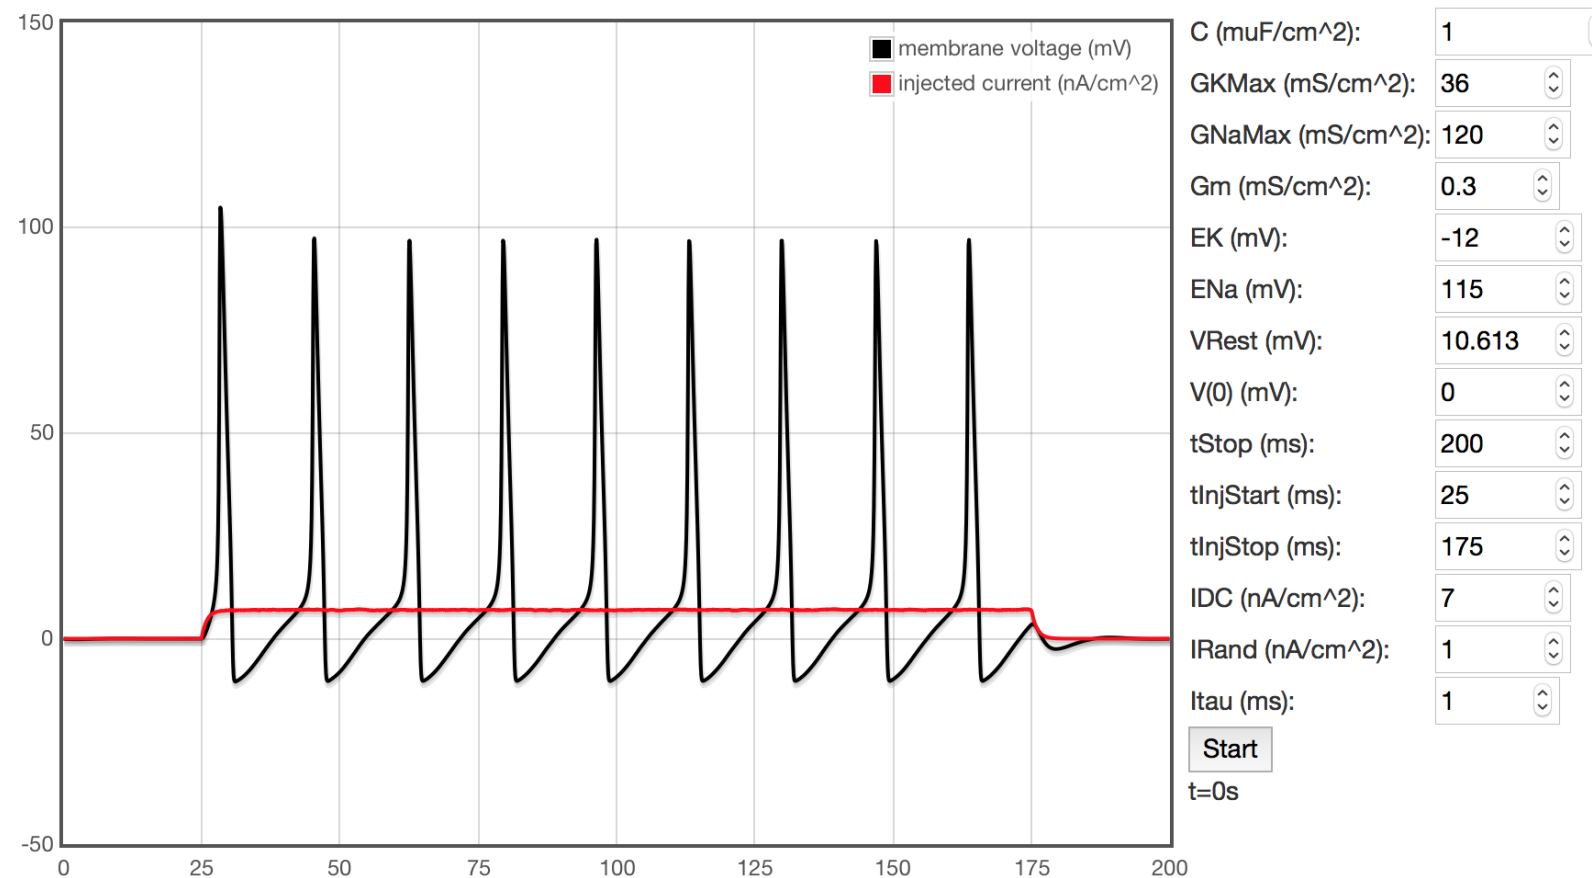


**neurons have a
nonlinear
response function**

**an action potential is generated
if net input is greater than
an activation threshold**

action potential simulator

Hodgkin-Huxley Simulation with Javascript




<http://myselfph.de/hodgkinHuxley.html>

IDC (input/injected current)

IRand (random variability in current)

try IDC=2.355 vs. IDC=2.356 with IRand=0




**an action potential is generated
if net input is greater than
an activation threshold**

action potential is nature's way of
making the axon conductive over
relatively long distances


ion channels can open or
close depending on the
local voltage gradient



(A)




passive propagation of sub-threshold potential

(B)



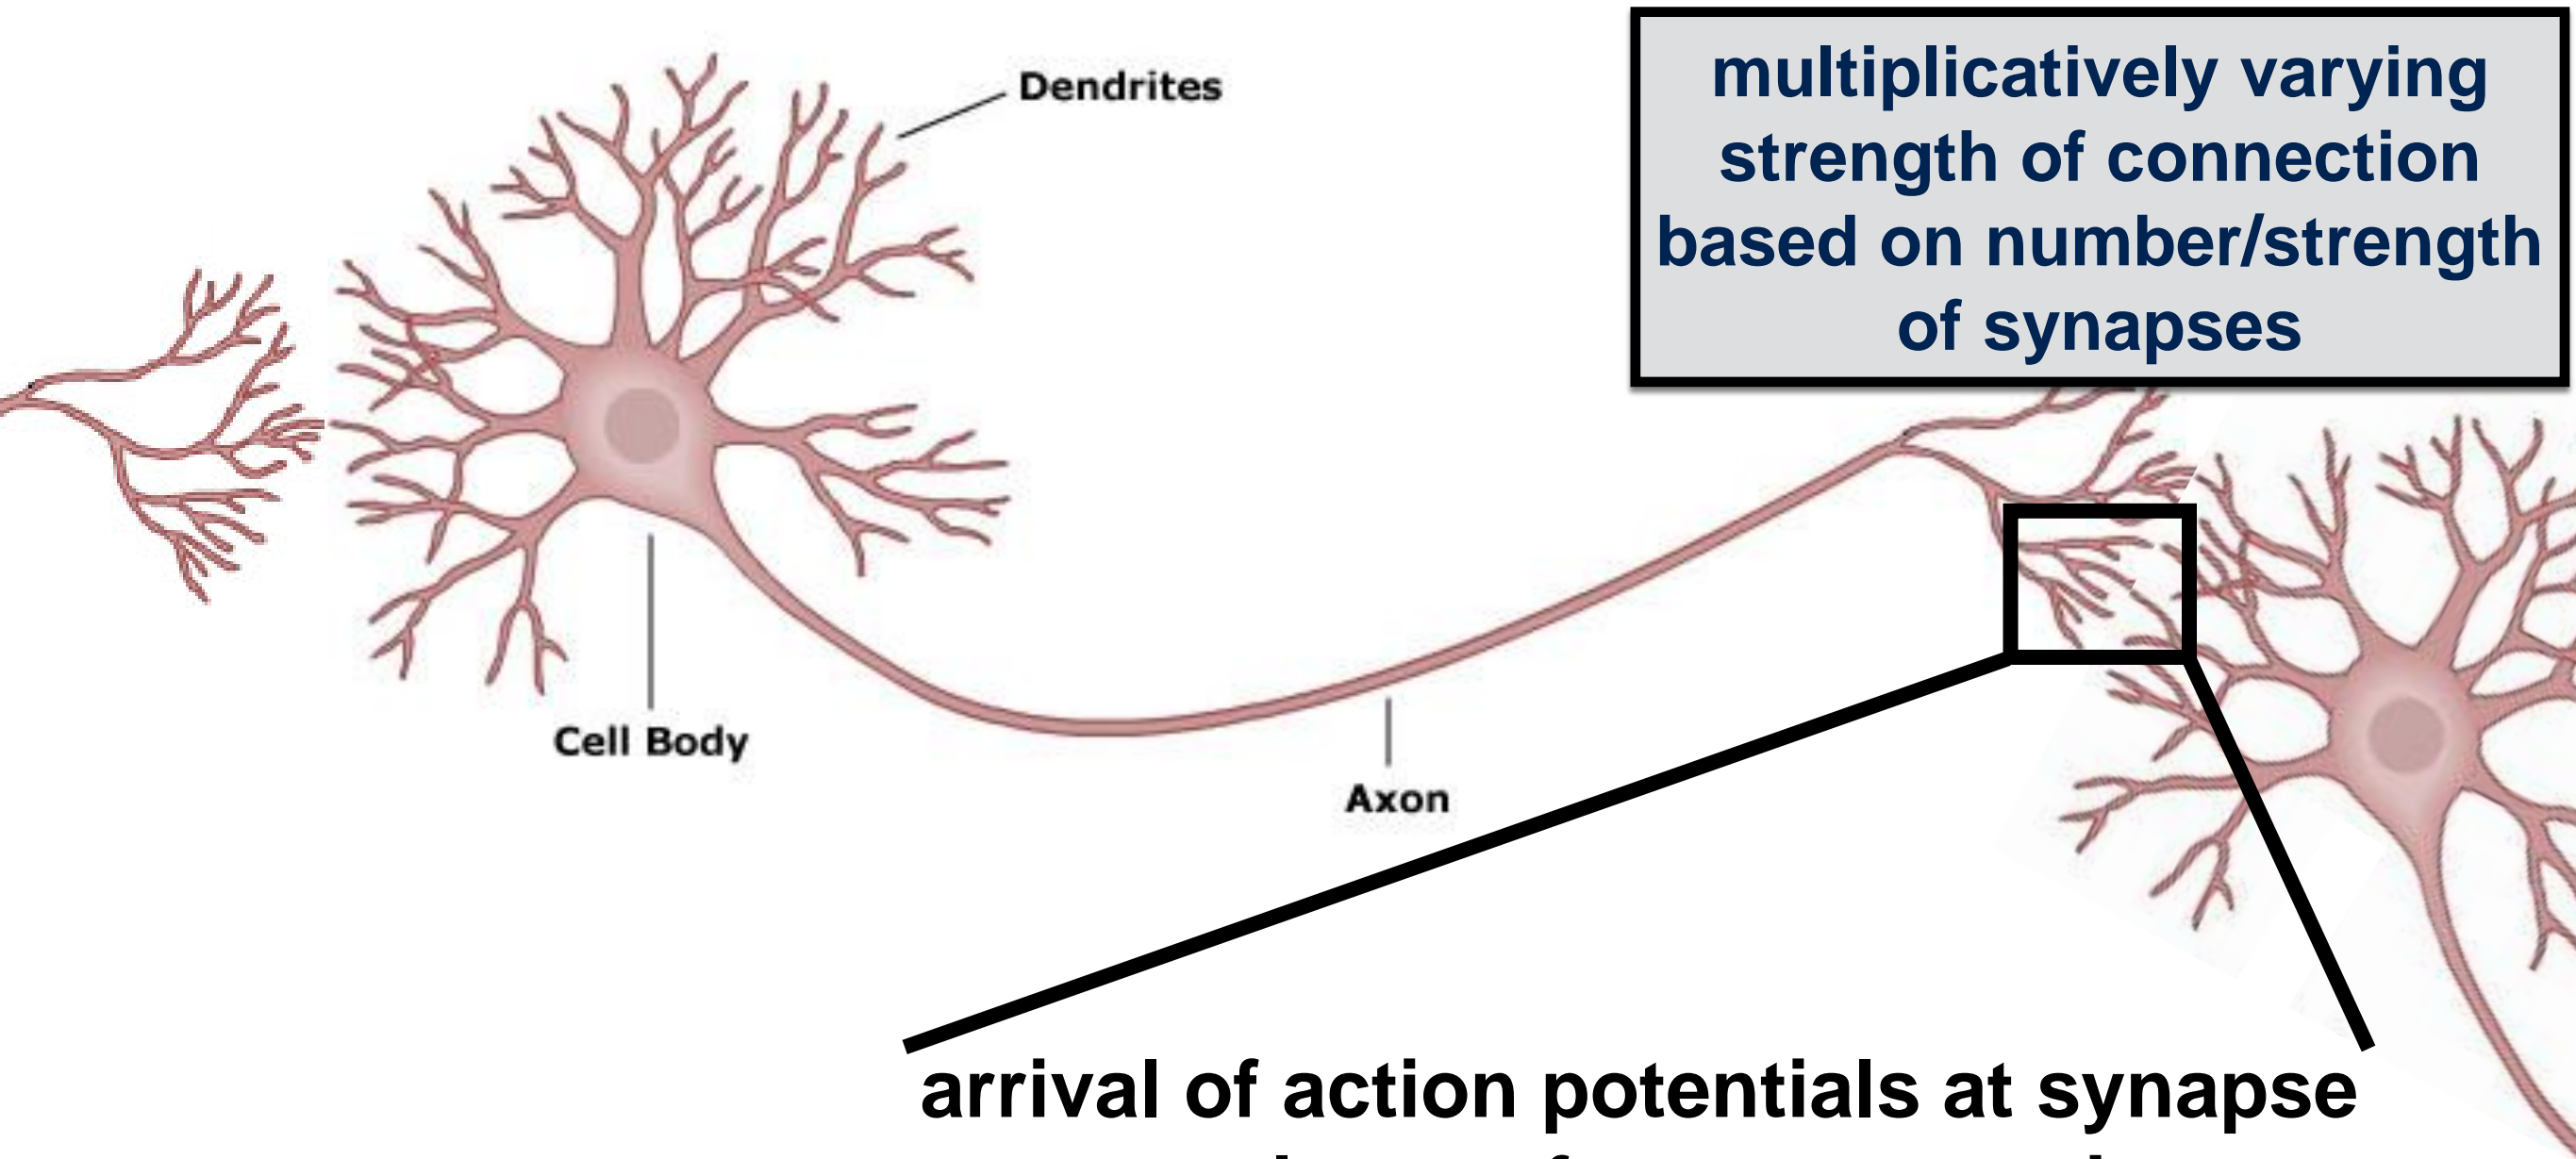




**arrival of action potentials at synapse
causes release of neurotransmitters to
post-synaptic neuron**

more and/or stronger synapses mean pre-
synaptic neuron has more influence over
post-synaptic neuron

synapses multiply effect of action potentials



**multiplicatively varying
strength of connection
based on number/strength
of synapses**

**arrival of action potentials at synapse
causes release of neurotransmitters to
post-synaptic neuron**


more and/or stronger synapses mean pre-
synaptic neuron has more influence over
post-synaptic neuron

synapses multiply effect of action potentials

Levels of Neuron Modeling

squid giant axon

- detailed model of action potentials (Hodgkin-Huxley model)



$$\frac{dV}{dt} = \frac{1}{C} (-I_{NA} - I_K - I_{leak} - I_{input})$$

$$\frac{dV}{dt} = \frac{1}{C} \left(-g_{NA} m^3 h (V - E_{NA}) - g_K n^4 (V - E_K) - g_{leak} (V - E_{leak}) + I_{input} \right)$$

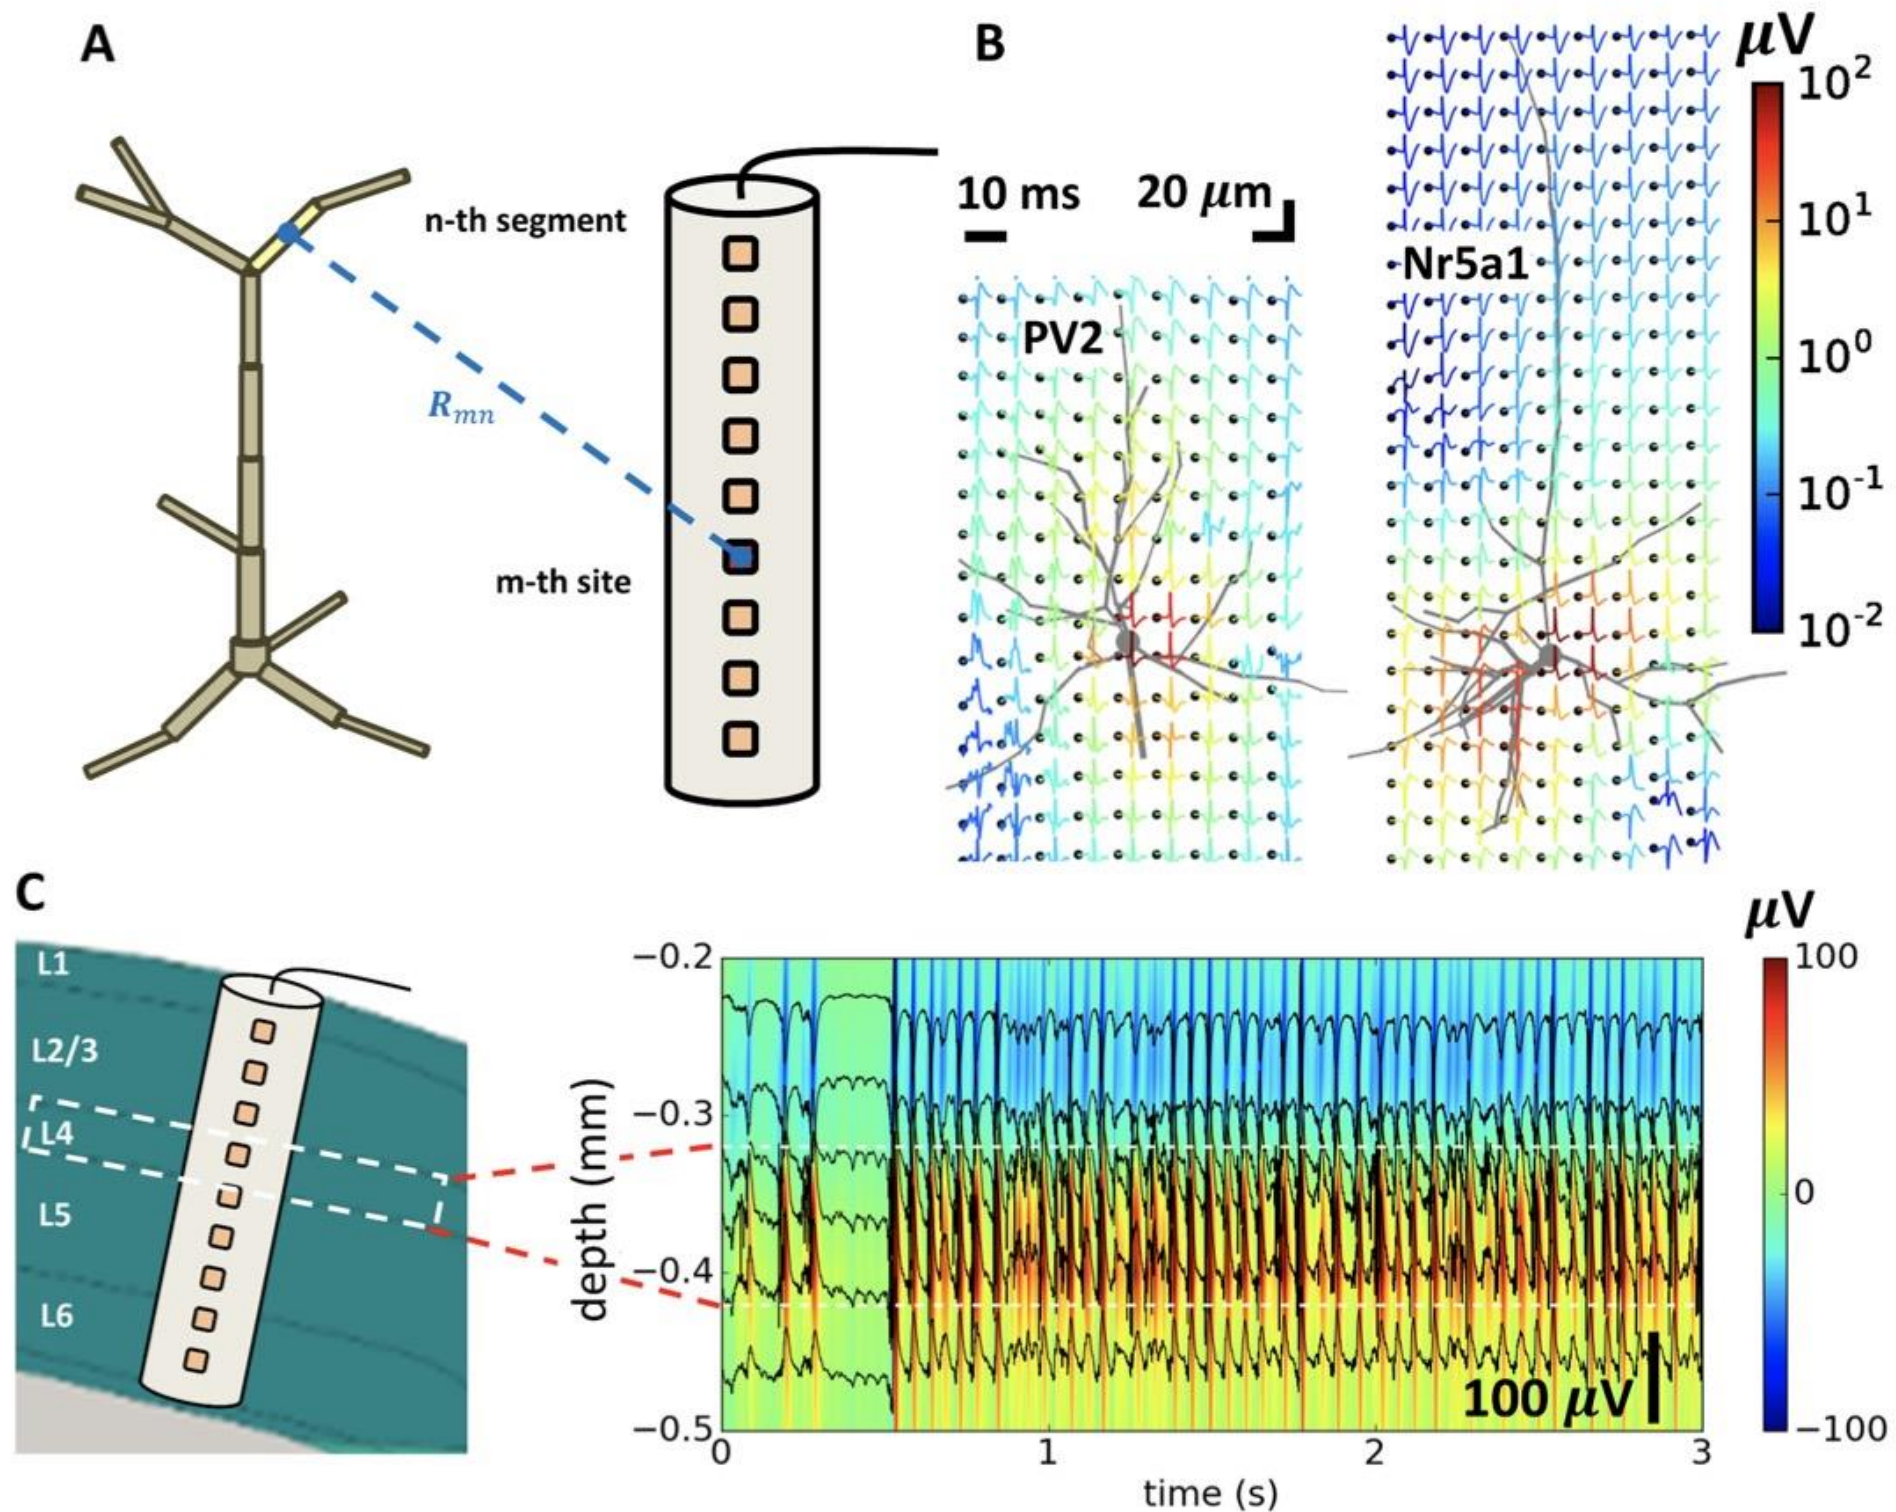
$$\frac{dm}{dt} = \frac{1}{\tau_m(V)} (-m + M(V))$$

$$\frac{dh}{dt} = \frac{1}{\tau_h(V)} (-h + H(V))$$

$$\frac{dn}{dt} = \frac{1}{\tau_n(V)} (-n + N(V))$$


Levels of Neuron Modeling

- more detailed model of neuron morphology (NEURON simulator)



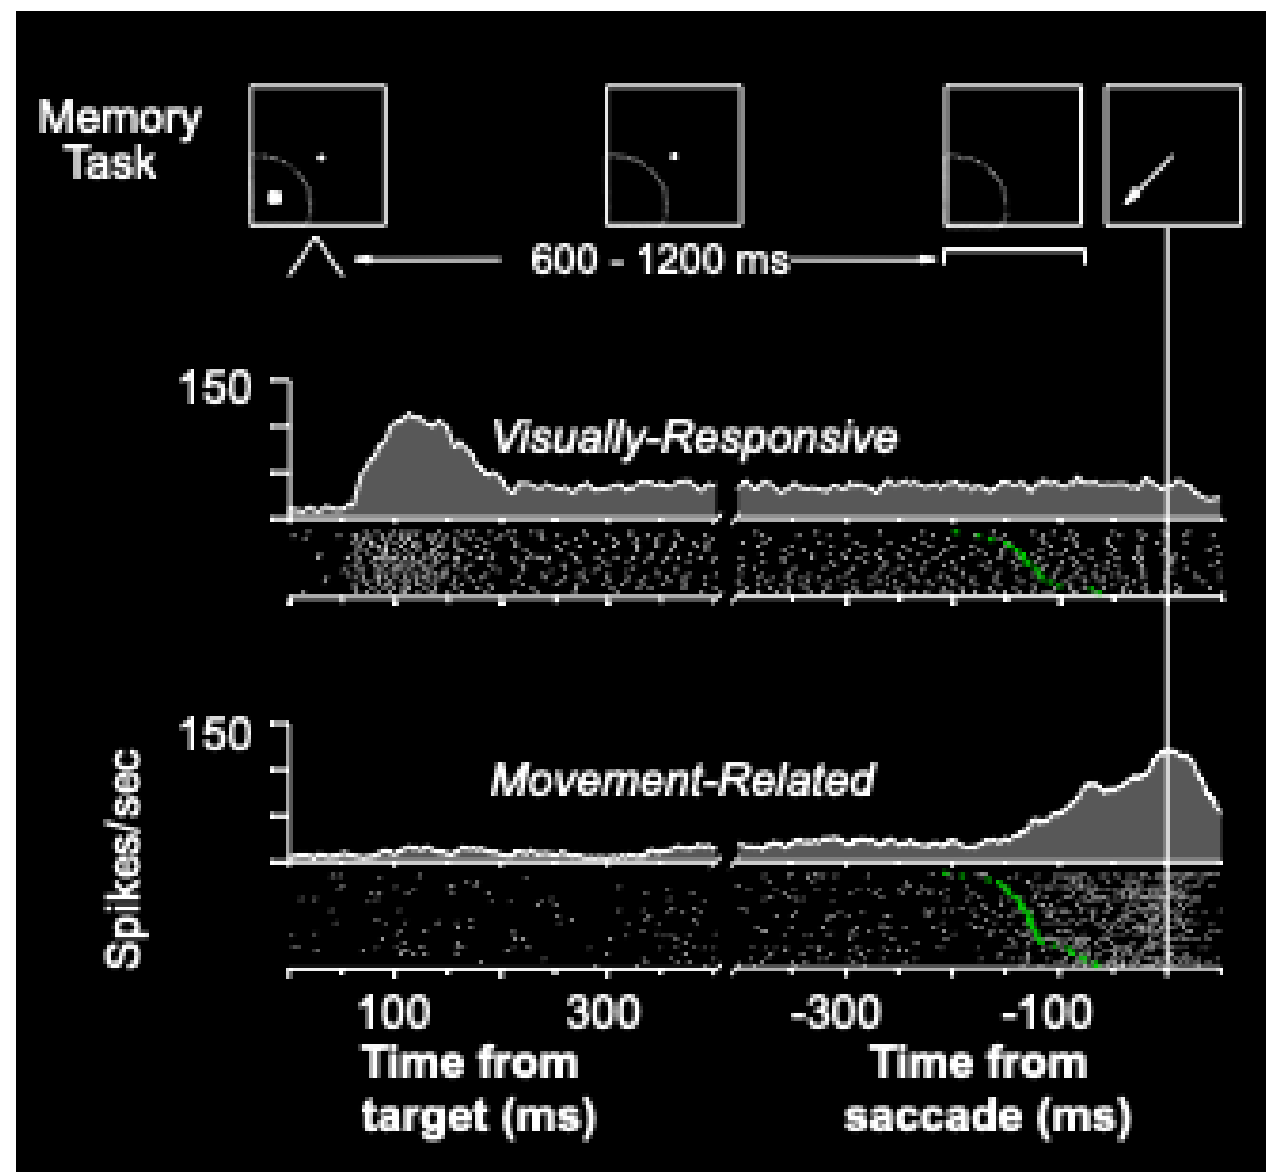
Levels of Neuron Modeling

- detailed model of action potentials (Hodgkin-Huxley model)
- more abstract integrate-and-fire (spiking) neuron




Levels of Neuron Modeling

- detailed model of action potentials (Hodgkin-Huxley model)
- more abstract integrate-and-fire (spiking) neuron
- rate-coded neuron



Memory Task



Levels of Neuron Modeling

- detailed model of action potentials (Hodgkin-Huxley model)
- more abstract integrate-and-fire (spiking) neuron
- rate-coded neuron
- activation (positive or negative) (collection of neurons)
what we'll be focusing on

Levels of Neuron Modeling

- detailed model of action potentials (Hodgkin-Huxley model)
- more abstract integrate-and-fire (spiking) neuron
- rate-coded neuron
- activation (positive or negative) (collection of neurons)



we'll start by considering neural network models that ignore spikes, and treat time in a “chunky” way

then consider models that learn, networks that change over time with experience



may consider models that allow for dynamics (recurrent networks, dynamical networks, spiking)

Introduction to Neural Networks



modeling an idealized neuron





Idealized Neuron



Idealized Neuron



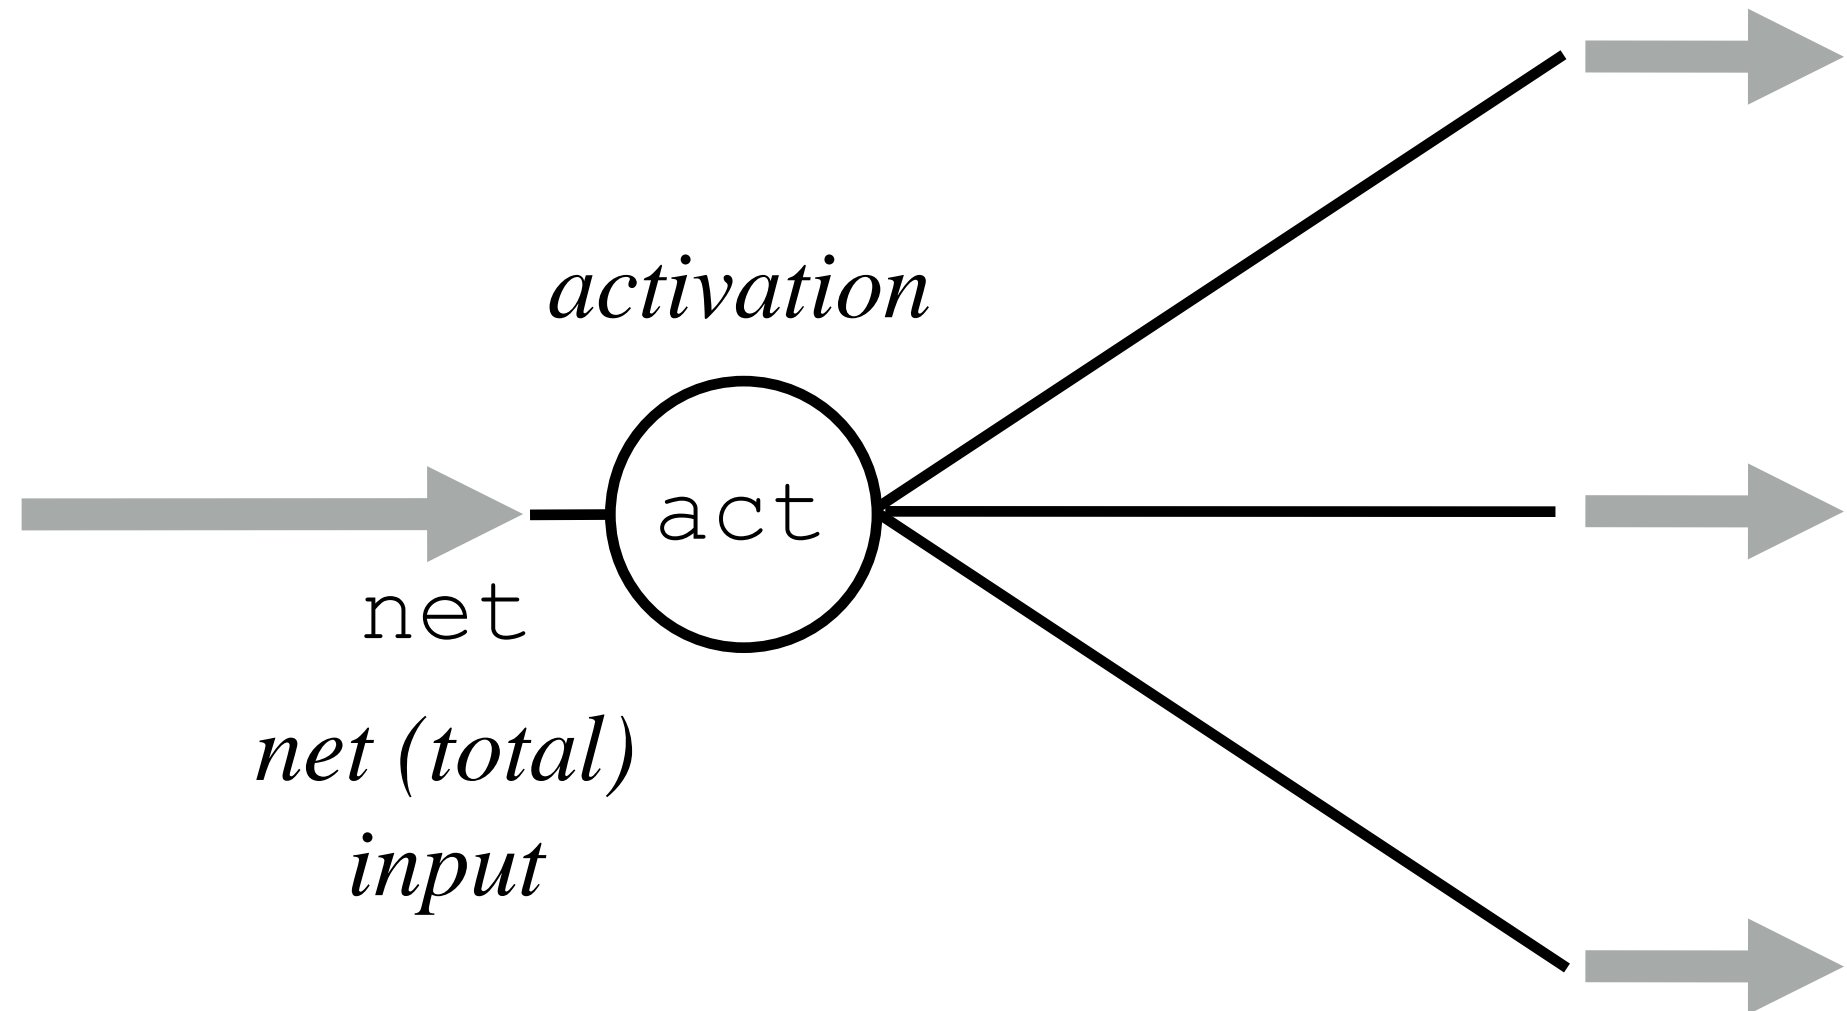
Idealized Neuron



Idealized Neuron

"unit" or "node"

how is the activation of a neuron related to its net input?



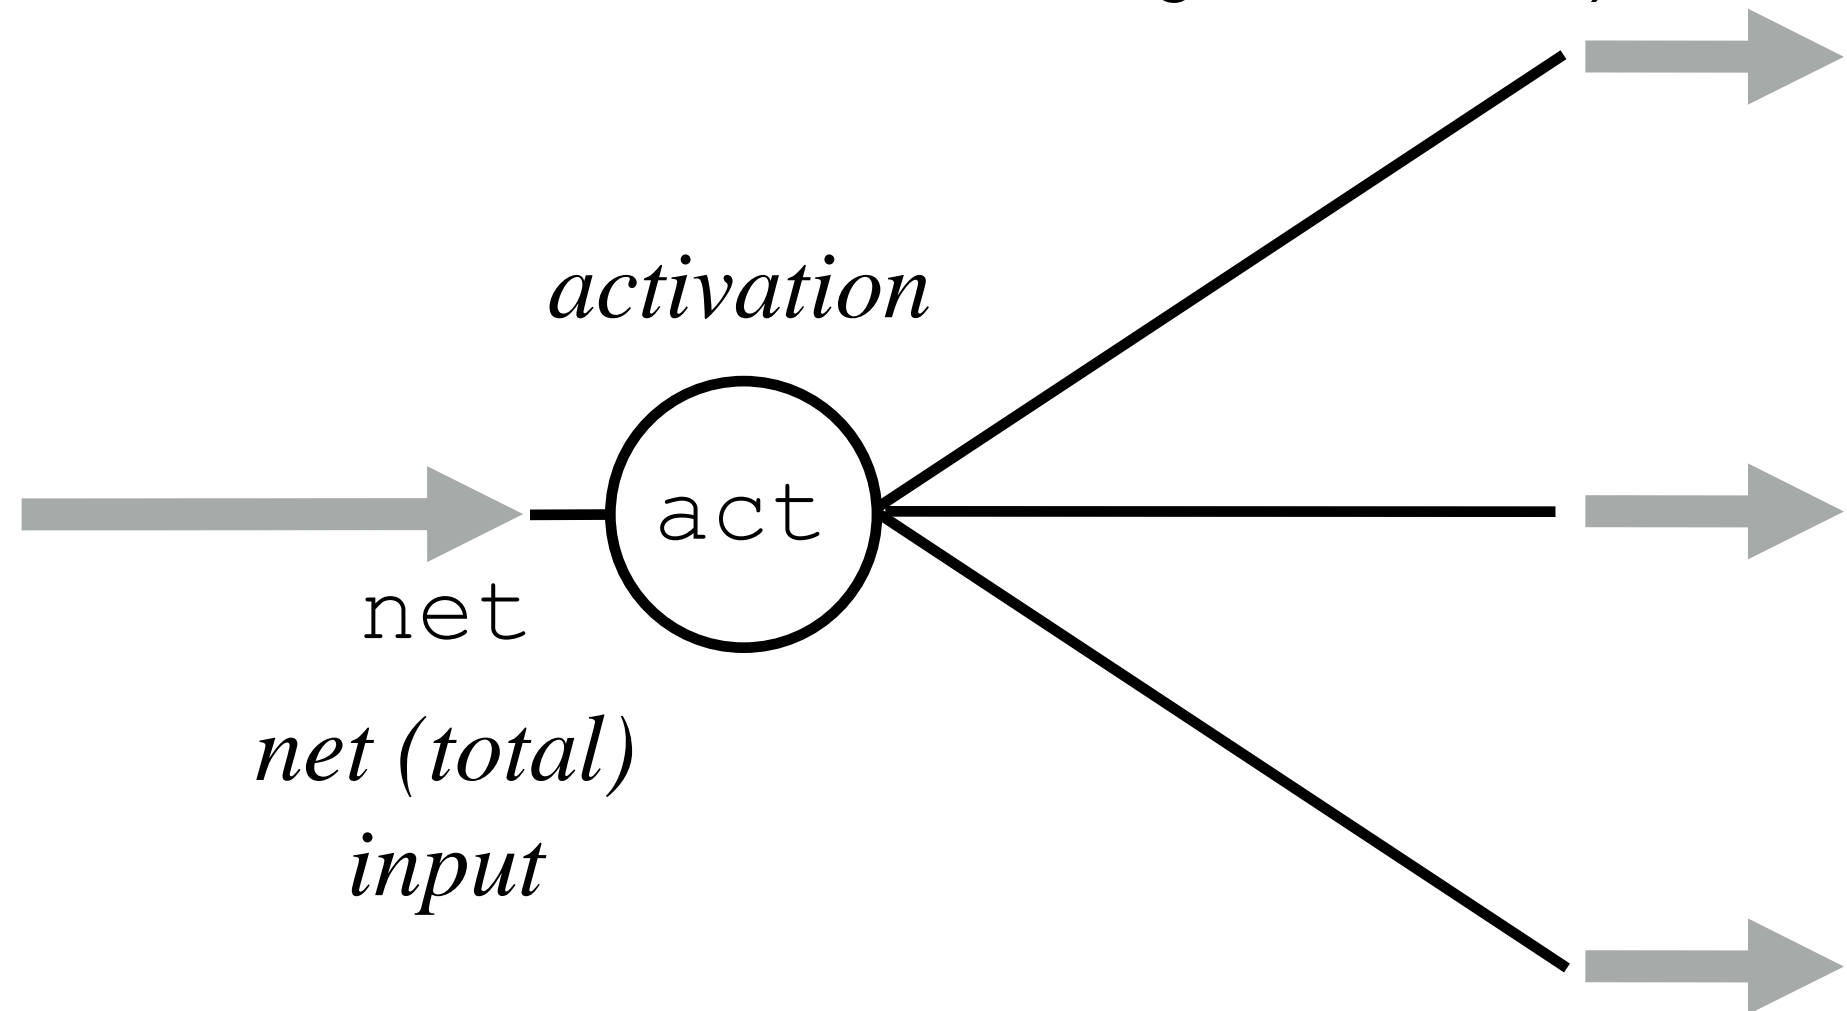
Idealized Neuron

"unit" or "node"

how is the activation of a neuron related to its net input?

*where "net input" comes from other neurons (and their synapses)
(where net input can be positive or negative)*


*where "activation" is like spike rate
(but some activation functions allow negative values)*



Idealized Neuron

*a variety of activation functions are used
to model neurons in neural networks*

*because there are always multiple neurons (units/nodes)
in a neural network we will commonly use subscripts*



Idealized Neuron

Keras, Tensorflow, and Python (plus toolkits)

Python

some programming will be
in base Python (with packages like
numpy, matplotlib, etc.)

Keras, Tensorflow, and Python (plus toolkits)

Python

some programming will be
in base Python (with packages like
numpy, matplotlib, etc.)


Keras

high-level Python package for
simulating neural networks

Tensorflow

high-level Python package for
simulating neural networks

CPUs, GPUs, TPUs





used as output activation
when modeling "regression"
problems

not used as activation
for "hidden" units
(we will discuss later)

linear activation function

$$a_j(n_j) = n_j$$





used as output activation
when modeling "regression"
problems

not used as activation
for "hidden" units
(we will discuss later)

linear activation function


```
import numpy as np
import tensorflow as tf
```

```
n = np.linspace(-6, 6, 1000)
a = tf.keras.activations.linear(n)
```


```
my_singleplot(n, a, 'linear (TF)')
```

```
a = tf.keras.activations.linear(n)
```

**neurons have a
nonlinear
response function**




sigmoid (logistic)




$$a_j(n_j) = \frac{1}{1 + \exp(-n_j)}$$

```
a = tf.keras.activations.sigmoid(n)
```



**neurons have a
nonlinear
response function**


tanh (hyperbolic tangent)




$$a_j(n_j) = \frac{\exp(n_j) - \exp(-n_j)}{\exp(n_j) + \exp(-n_j)}$$


```
a = tf.keras.activations.tanh(n)
```


where does the net input come from?




where does the net input come from?






where does the net input come from?




**inputs can be
positive (+) or
negative(-)**



while individual neurons are either excitatory or inhibitory, inhibitory interneurons can make an excitatory neuron inhibit indirectly

inputs can be positive (+) or negative(-)

so for simplicity, we allow the same unit to be either excitatory or inhibitory




Two different implementations

Whether or not you allow the same unit to be excitatory and inhibitory.

This may not affect the kinds of computations a network can perform.


activation of
pre-synaptic
neurons

Note: because neural networks have many
nodes and many connections, we will
have to make extensive use of subscripts.




activation of
pre-synaptic
neurons

and sometimes superscripts as well ...
if unsure about notation, please ask ...
also note that notation may not always
be constant across things we read




inputs clamped
onto network



$$w_{ij}$$

weight of synaptic connections


subscripts on weights will generally be
from (which unit) followed by to (which unit)



$$w_{ij}$$

weight of synaptic connections
can vary in strength (and sign)


subscripts on weights will generally be
from (which unit) followed by to (which unit)



$$a_i w_{ij}$$

impact of one pre-synaptic neuron

**multiplicatively varying
strength of connection
based on number/strength
of synapses**




$$a_i w_{ij}$$

why is it multiplicative
and not additive?

*multiplication is more
than just repeated addition*


**multiplicatively varying
strength of connection
based on number/strength
of synapses**



$$a_i w_{ij}$$

if the pre-synaptic neuron
is inactive, it does not matter
how strong the connection


**multiplicatively varying
strength of connection
based on number/strength
of synapses**



$$a_i w_{ij}$$

if the connection is weak or absent,
it does not matter if pre-synaptic
neuron is active or not


**multiplicatively varying
strength of connection
based on number/strength
of synapses**



net input sums the weighted inputs


$$n_j = \sum_i a_i w_{ij}$$

**inputs integrate
at the cell body
- they are added
together**



let's take a look at this equation


$$n_j = \sum_i a_i w_{ij}$$



let's take a look at this equation


$$n_j = \sum_i a_i w_{ij}$$

$$n_j = \sum_{i=0}^{M-1} a_i w_{ij}$$



let's take a look at this equation

$$n_j = \sum_i a_i w_{ij}$$



$$n_j = \sum_{i=0}^{M-1} a_i w_{ij} = a_0 w_{0j} + \dots + a_i w_{ij} + \dots + a_{M-1} w_{M-1,j}$$

$$a = [a_0, \dots, a_i, \dots, a_{M-1}]$$

vector of activations of
the input nodes


$$w_j = [w_{0j}, \dots, w_{ij}, \dots, w_{M-1,j}]$$

vector of weights on
connections to node j

see `Week3b.ipynb`

let's take a look at this equation

$$n_j = \sum_i a_i w_{ij}$$




$$n_j = \sum_{i=0}^{M-1} a_i w_{ij} = \boxed{a_0} w_{0j} + \dots + \boxed{a_i} w_{ij} + \dots + \boxed{a_{M-1}} w_{M-1,j}$$

$$a = [a_0, \dots, a_i, \dots, a_{M-1}]$$

$$w_j = [w_{0j}, \dots, w_{ij}, \dots, w_{M-1,j}]$$

let's take a look at this equation

$$n_j = \sum_i a_i w_{ij}$$



$$n_j = \sum_{i=0}^{M-1} a_i w_{ij} = a_0 w_{0j} + \dots + a_i w_{ij} + \dots + a_{M-1} w_{M-1,j}$$

$$a = [a_0, \dots, a_i, \dots, a_{M-1}]$$

$$w_j = [w_{0j}, \dots, w_{ij}, \dots, w_{M-1,j}]$$


let's take a look at this equation

$$n_j = \sum_i a_i w_{ij}$$

$$a = [a_0, \dots, a_i, \dots, a_{M-1}]$$

$$w_j = [w_{0j}, \dots, w_{ij}, \dots, w_{M-1,j}]$$

```
net = 0.0
for i in range(len(a)) :
    net += a[i]*w[j][i]
```



vector of activations of
the input nodes

vector of weights on
connections to node j

using for loops


let's take a look at this equation

$$n_j = \sum_i a_i w_{ij}$$

$$a = [a_0, \dots, a_i, \dots, a_{M-1}]$$

$$w_j = [w_{0j}, \dots, w_{ij}, \dots, w_{M-1,j}]$$

```
net = np.sum(a*wj)
```



vector of activations of
the input nodes

vector of weights on
connections to node j

vectorized

let's take a look at this equation

$$n_j = \sum_i a_i w_{ij}$$


$$a = [a_0, \dots, a_i, \dots, a_{M-1}]$$

$$w_j = [w_{0j}, \dots, w_{ij}, \dots, w_{M-1,j}]$$

numpy array (from element-wise multiplication)

```
net = a*wj
net = np.sum(net)
```

float (sum of elements of numpy array)



vector of activations of
the input nodes

vector of weights on
connections to node j

vectorized


let's take a look at this equation

$$n_j = \sum_i a_i w_{ij}$$

$$a = [a_0, \dots, a_i, \dots, a_{M-1}]$$

$$w_j = [w_{0j}, \dots, w_{ij}, \dots, w_{M-1,j}]$$

```
net = np.dot(a, wj)
```



vector of activations of
the input nodes

vector of weights on
connections to node j

as vector dot product

let's take a look at this equation


$$n_j = \sum_i a_i w_{ij} = \boxed{a \cdot w_j}$$

$$a = [a_0, \dots, a_i, \dots, a_{M-1}]$$

$$w_j = [w_{0j}, \dots, w_{ij}, \dots, w_{M-1,j}]$$

```
net = np.dot(a, wj)
```

linear algebra operation on vectors



vector of activations of
the input nodes

vector of weights on
connections to node j


as vector dot product

dot product

computes a “similarity” between two vectors

```
a = np.array([1, 3])
```

```
wj = np.array([2, 1])
```



dot product

angle between two vectors ...


$$\cos(\theta) = \frac{a \cdot w_j}{\|a\| \|w\|}$$

← dot product

← norms (lengths)

$$\|a\| = \sqrt{a \cdot a}$$


just the Pythagorean theorem



dot product


angle between two vectors ...

$$\cos(\theta) = \frac{a \cdot w_j}{\|a\| \|w\|}$$




dot product

```
costheta = np.dot(a,wj) /  
            (np.sqrt(np.dot(a,a)) *  
             np.sqrt(np.dot(wj,wj)))  
theta = np.arccos(costheta)
```



dot product


```
costheta = np.dot(a,wj) /  
            (np.linalg.norm(a) *  
             np.linalg.norm(wj))  
theta = np.arccos(costheta)
```



dot product

another way to write this (rearrange terms)

$$\cos(\theta) = \frac{a \cdot w_j}{\|a\| \|w\|}$$



dot product

dot product equals product of the lengths and the angle

$$a \cdot w_j = \|a\| \|w_j\| \cos(\theta)$$


dot
product

length
 a

length
 w_j

cos
angle

what happens if vectors are orthogonal?



dot product

dot product equals product of the lengths and the angle


$$a \cdot w_j = \|a\| \|w_j\| \cos(\theta)$$

dot
product

length
 a

length
 w_j

cos
angle



what happens if vectors are orthogonal?

$\theta = 90$ degrees
so $\cos(\theta) = 0.0$

dot product

dot product equals product of the lengths and the angle


$$a \cdot w_j = \|a\| \|w\| \cos(\theta)$$

dot
product

length
 a

length
 w_j

cos
angle



what happens if vectors are orthogonal?

$\theta = 90$ degrees
so $\cos(\theta) = 0.0$

all Python code uses radians (not degrees)
 $\theta = \pi/2$ radians = 90 degrees
so $\cos(\theta) = 0.0$


so dot product is zero

dot product


dot product is general

dot product lets you compute the angle
(a measure of similarity) between vectors
that are 2D, 10D, or 100D

$$\cos(\theta) = \frac{a \cdot w_j}{\|a\| \|w\|}$$




layers of a neural network



Idealized Neuron


fully-interconnected two-layer network

this is how some readings
will describe networks




fully-interconnected
~~**two-layer network**~~
single-layer network

this is how **keras**
will define networks




fully-interconnected three-layer network

this is how some readings
will describe networks




fully-interconnected
~~**three-layer network**~~
two-layer network

this is how **keras**
will define networks



bias term in net input



activation function

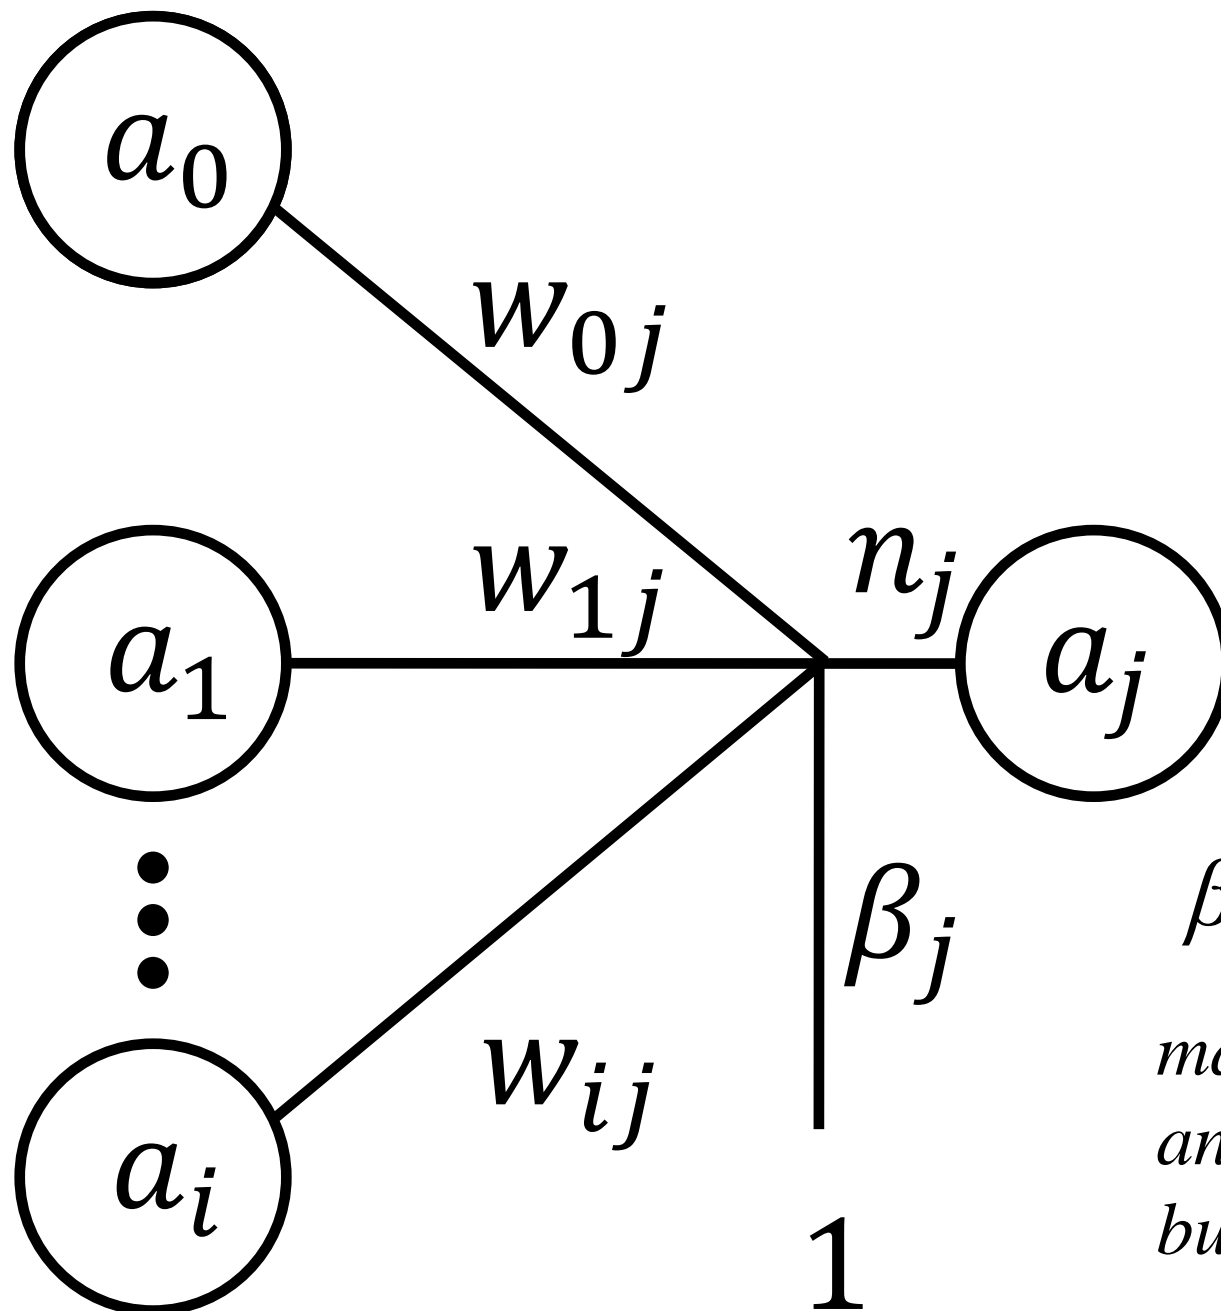
$$n_j = \sum_i a_i w_{ij}$$

$$a_j = f(n_j)$$

bias term in net input

$$n_j = \sum_i a_i w_{ij} + \beta$$

$$a_j = \frac{1}{1 + \exp(-n_j)}$$



β is the bias of a unit

*mathematically, it can be treated like
another weight contributing to net input,
but it is kept separate in keras*